

# Hadoop Distributed File System

Big Data Management

# Knowledge objectives

1. Recognize the need of persistent storage
2. Enumerate the design goals of GFS
3. Explain the structural components of HDFS
4. Name three file formats in HDFS and explain their differences
5. Recognize the importance of choosing the file format depending on workload
6. Explain the actions of the coordinator node in front of chunkserver failure
7. Explain a mechanism to avoid overloading the master node in HDFS
8. Explain how data is partitioned and replicated in HDFS
9. Recognize the relevance of sequential read

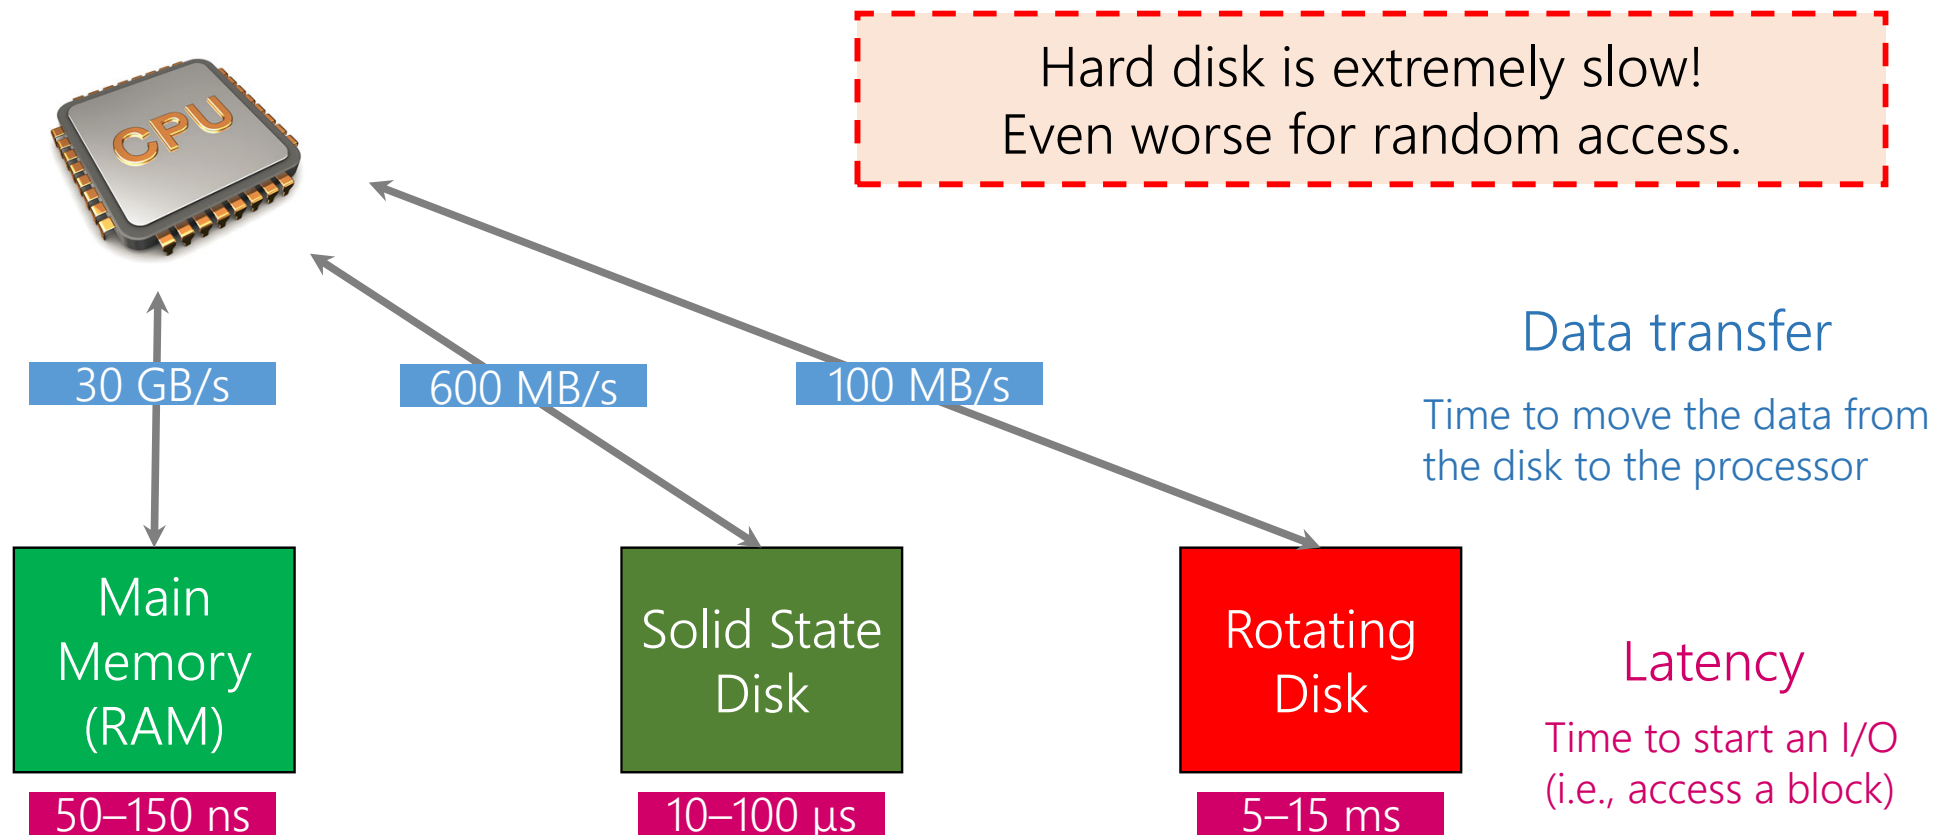
# Understanding objectives

1. Choose the format for an HDFS file based on heuristics
2. Estimate the data retrieved by scan, projection and selection operations in SequenceFile, Avro and Parquet

# (Distributed) File Systems

Google File System

# Time to bring data (approximations)



# Reasons to keep using HDDs



Main  
Memory

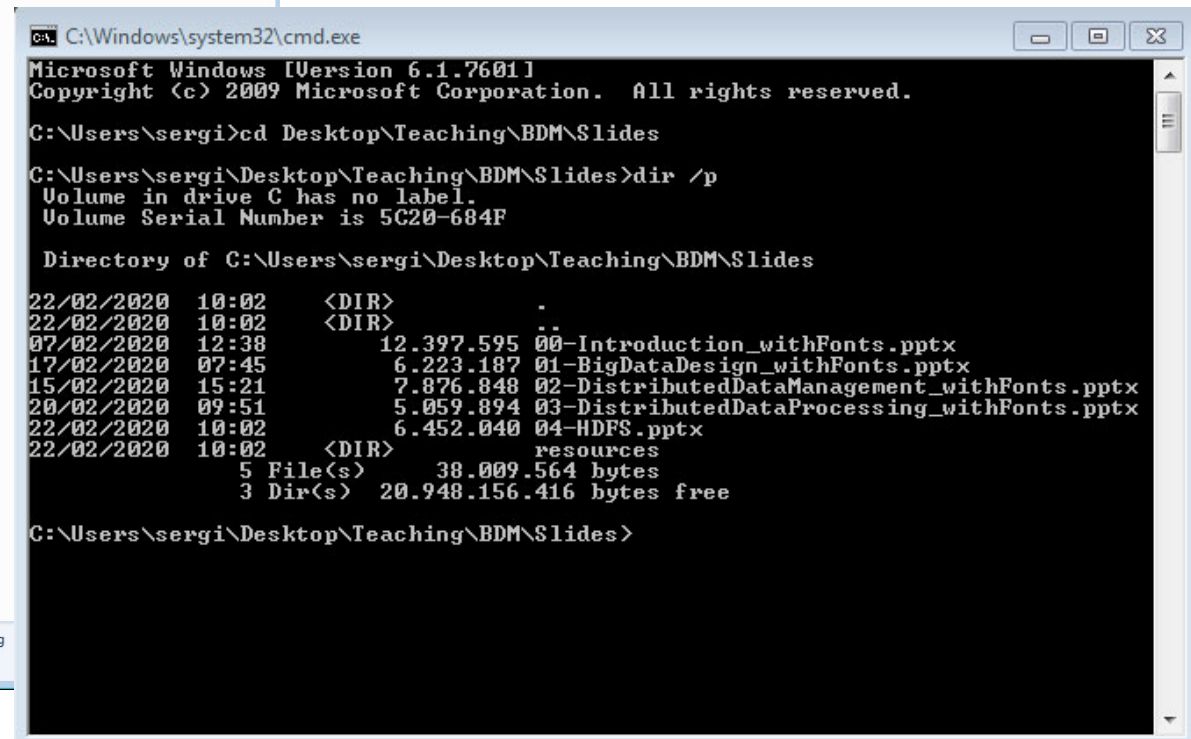
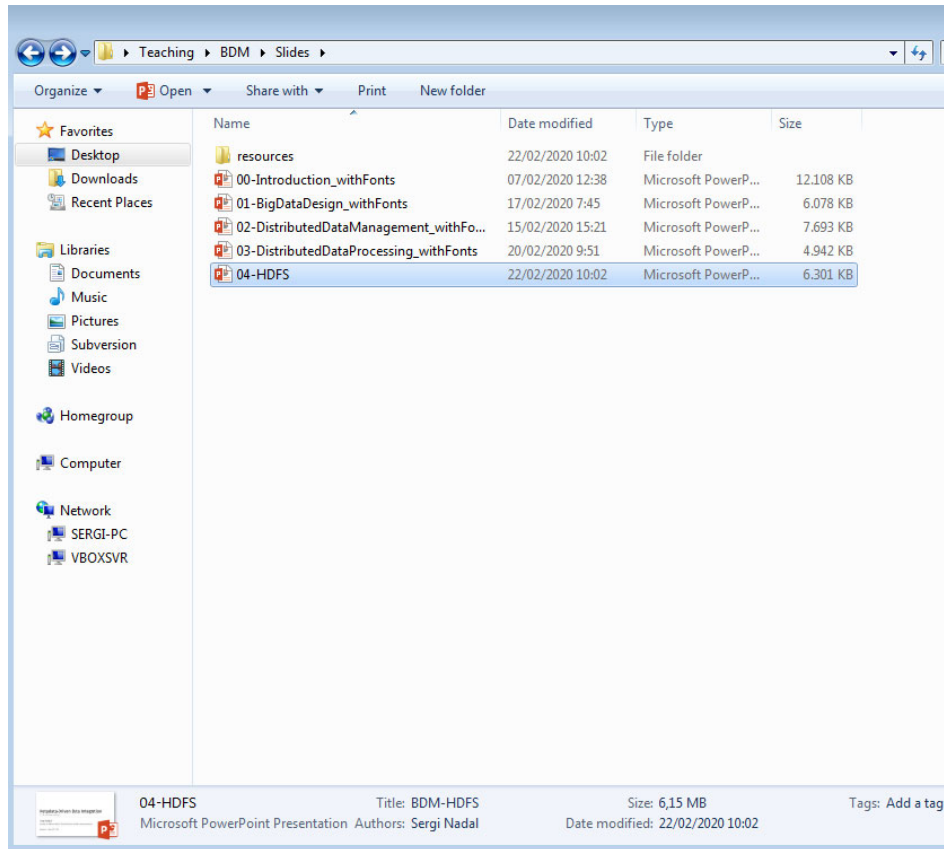
Faster ✓

✓ Cheaper  
✓ Persistent



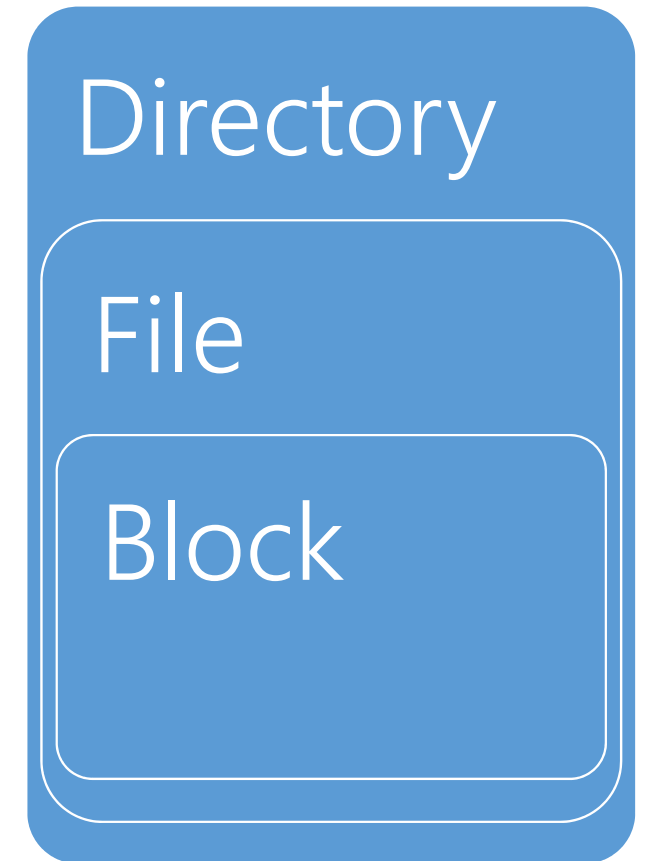
Hard  
Drive Disk

# What is a file system?



# Functionalities provided by a FS

- Creates a hierarchical structure of data
  - Splits and stores data into files and blocks
- Provides interfaces to read/write/delete
- Maintains directories/files metadata
  - Size, date of creation, permissions, ...





# Distributed File Systems

- Same requirements, different setting
  1. Files are huge for traditional standards
  2. Most files are updated by appending data rather than overwriting
    - Write Once and Read Many times (WORM)
  3. Component failures are the norm rather than the exception
- Google File System (GFS)
  - The first large-scale distributed file system
  - Capacity of a GFS cluster

Capacity	Nodes	Clients	Files
10 PB	10.000	100.000	100.000.000

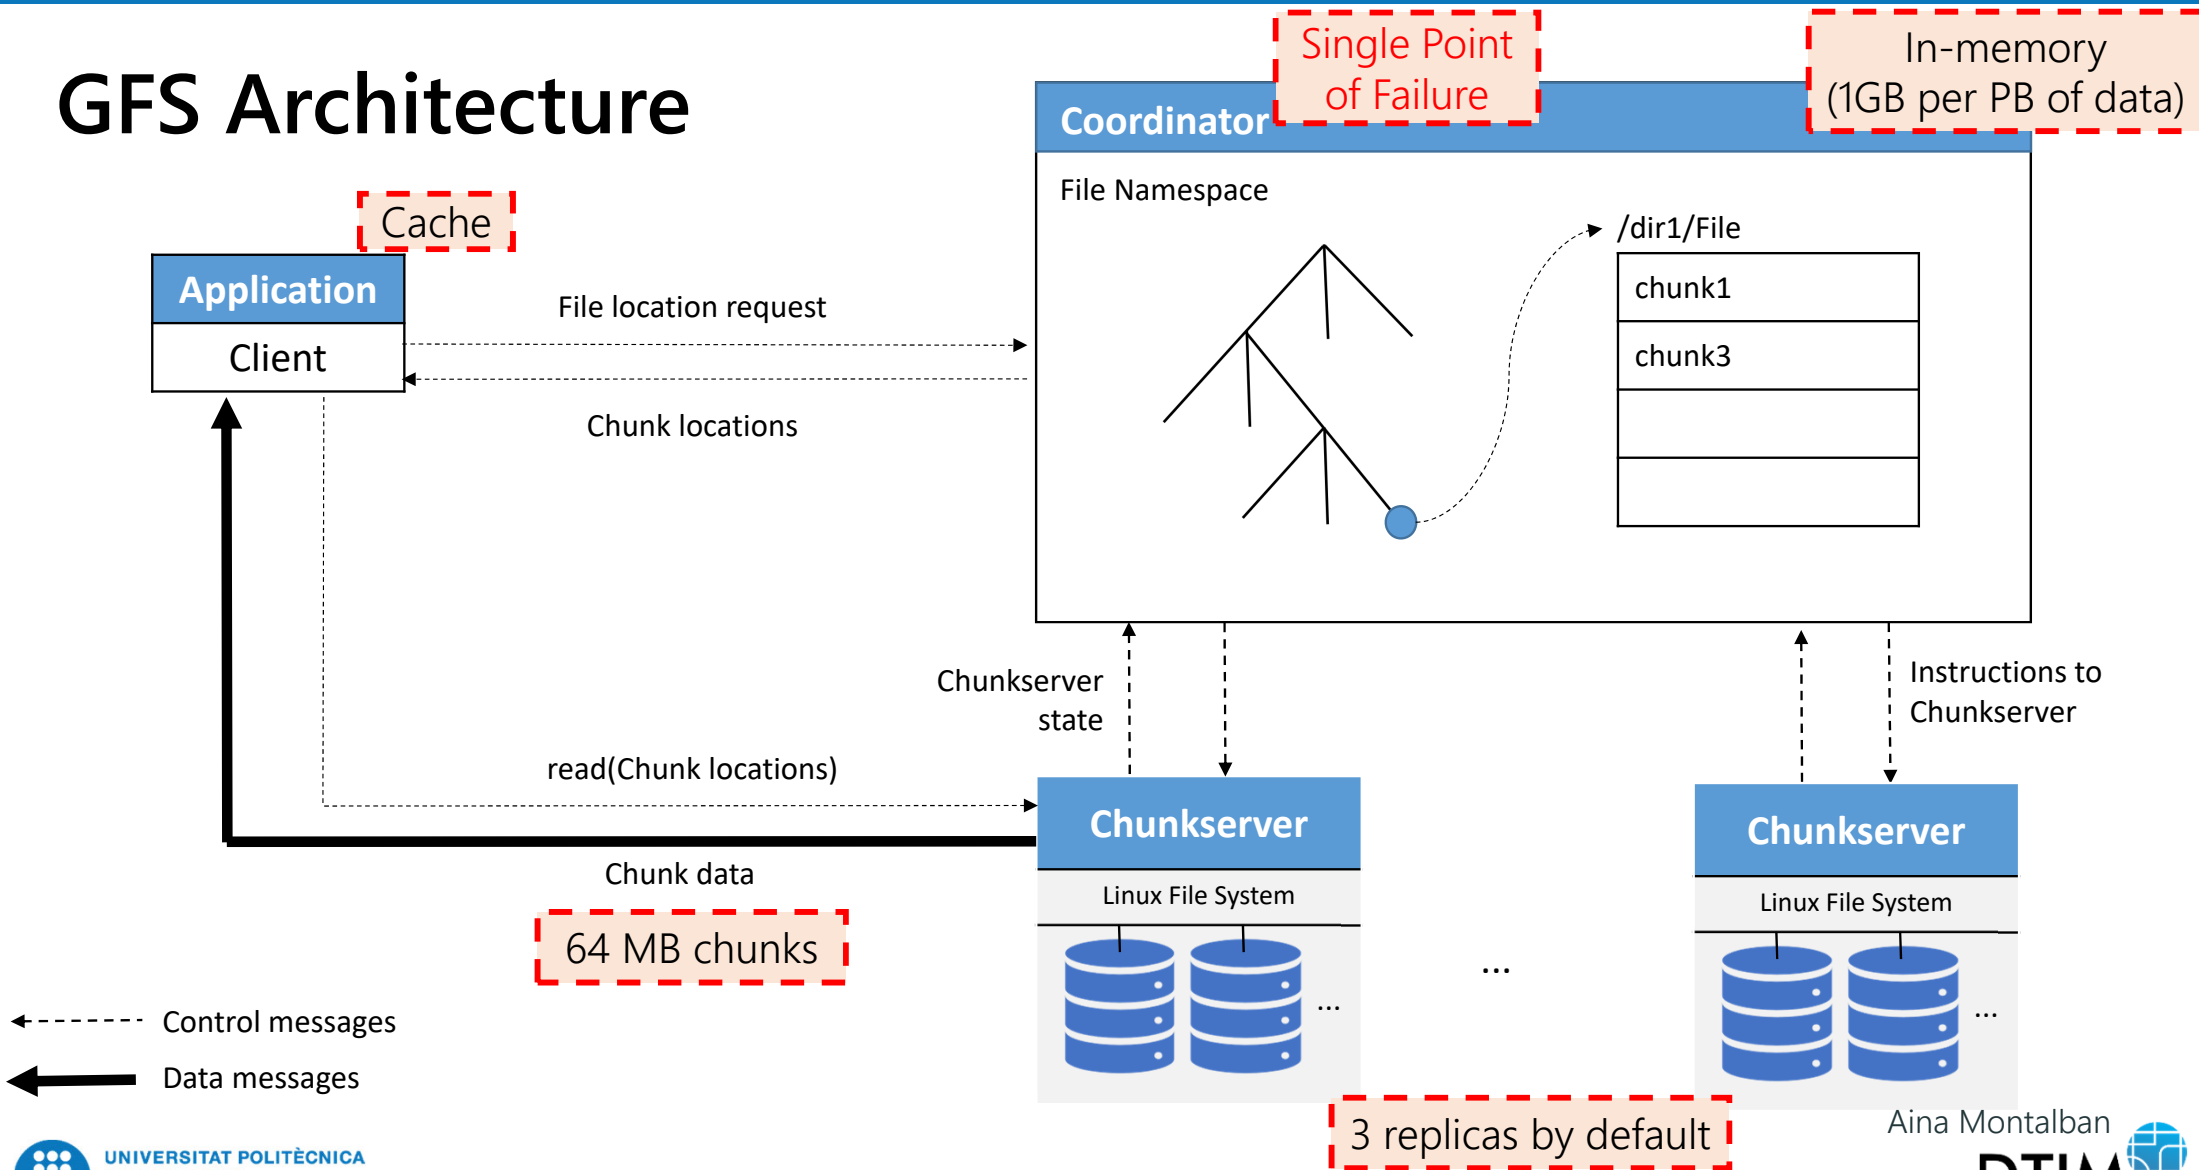
# Design goals of GFS

- Efficient management of files
  - Optimized for very large files (GBs to TBs)
- Efficiently append data to the end of files
  - Allow concurrency
- Tolerance to failures
  - Clusters are composed of many inexpensive machines that fail often
    - Failure probability (2-3/1.000 per day)
- Optimize sequential scans
  - Overcome high latency of HDDs (5-15ms) compared to main memory (50-150ns)

# GFS Architecture

There is a single coordinator, many chunk servers and many clients. Since there is only one coordinator, there might be a bottleneck. So, it needs to be very efficient. So, all the metadata is stored in memory. It keeps around 1 GB of metadata per PB of data. Still, it can fail. So, we need to reduce the risk of failure. This is done by caching the location of the chunk in the client side. If a client requests the same file twice, the request does not need to be made both times to the coordinator.

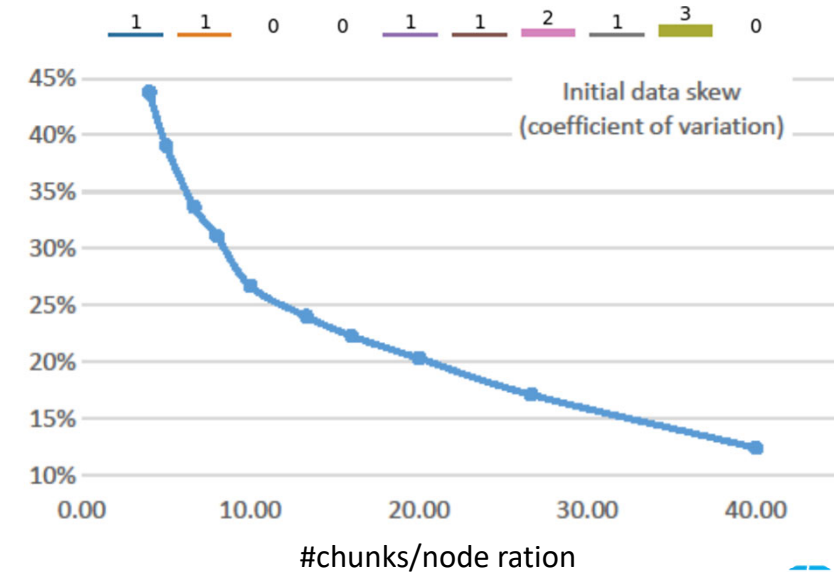
# GFS Architecture



# Other features

- Rebalance
  - Avoids skewness in the distribution of chunks
- Deletion
  - Moves a file to the trash (hidden)
    - Kept for 6h
  - *expunge* to force the trash to be emptied
- Management of stale replicas
  - Coordinator maintains versioning information about all chunks

Mean: 1 --- Stdev: 0.94 --- CoefOfVar: 0.94



# (Distributed) Data Design

Challenge I

# Storage layouts

“Jack of all trades, master of none”

- Different workloads require different layout
  - Horizontal
    - For scan-based workloads
  - Vertical
    - For projection-based workloads (reads a subset of columns)
  - Hybrid
    - For projection- and predicate-based workloads (reads a subset of columns or rows)

# Horizontal layout – Sequence File

- Records of binary key-value pairs

A	B	C	D
101	201	301	401
102	202	302	402
103	203	303	403

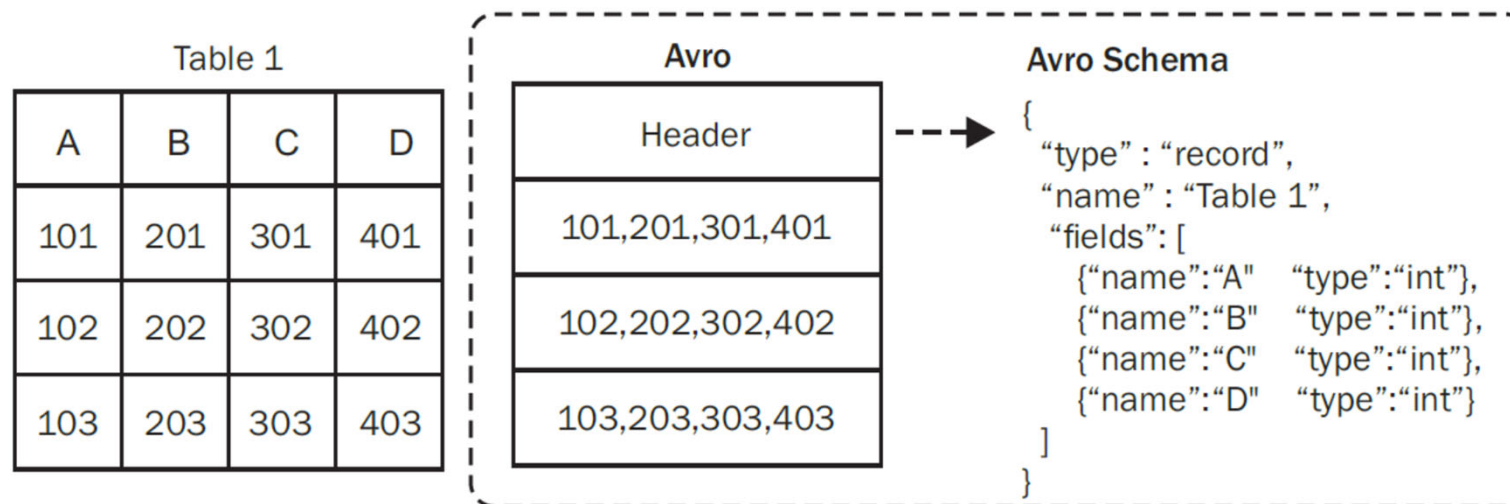
Header
Key : 101 Value : 201,301,401
Key : 102 Value : 202,302,402
Key : 103 Value : 203,303,403

- Compression
  - Uncompressed
  - Record-compressed
  - Block-compressed
    - “block” is the compression unit – a block of records (not a chunk)
      - 1 MB default



# Horizontal layout - Avro

- Binary encoding of (compressed) rows
- The header contains a schema encoded in JSON

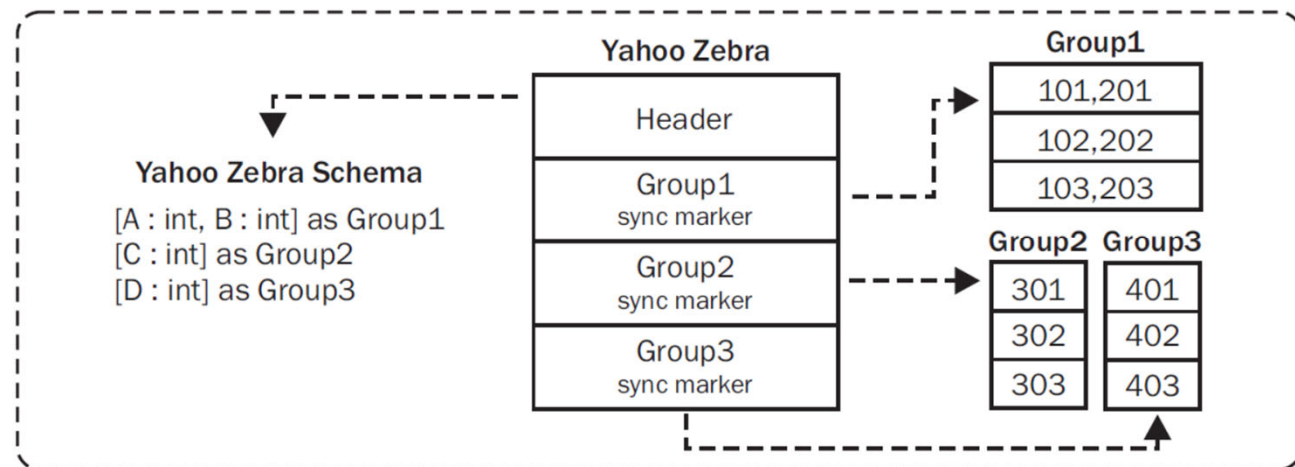


# Vertical layout - Zebra

- The header contains the definition of groups
  - Each group contains a set of columns
  - Widely benefits from compression
- Not really used in practice

Table 1

A	B	C	D
101	201	301	401
102	202	302	402
103	203	303	403

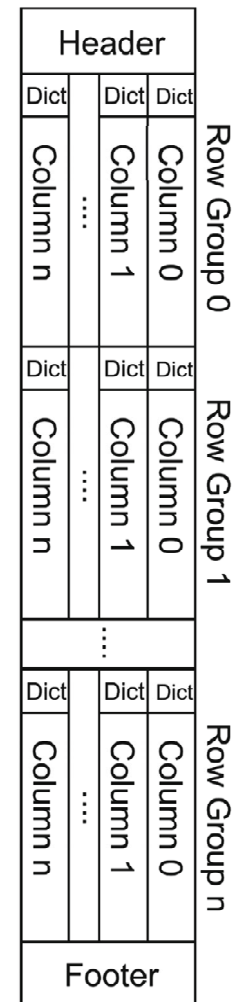
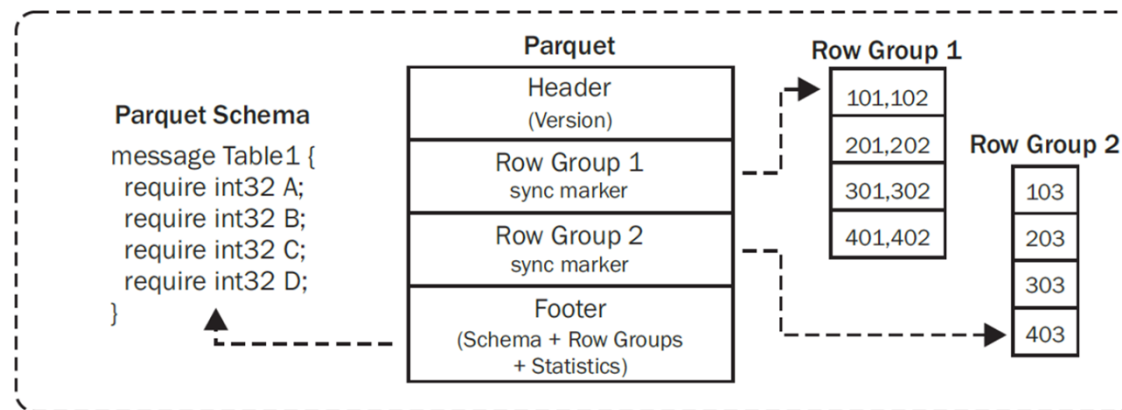


# Hybrid layout - Parquet

- Row groups (RG) - horizontal partitions
  - Data vertically partitioned within RGs
- Statistics per row group (aid filtering)
  - E.g., min-max

Table 1

A	B	C	D
101	201	301	401
102	202	302	402
103	203	303	403



# Parquet encoding definitions

- Plain
- Valid for any data type
  - Dictionary encoding
  - Run-length encoding with dictionary
- Specific for some kind of data types
  - Delta encoding
  - Delta byte array / Delta-length byte array
  - Byte Stream Split

# Comparison of data formats

Features	Horizontal		Vertical	Hybrid	
	Sequence Files	Avro	Yahoo Zebra	ORC	Parquet
Schema	No	Yes	Yes	Yes	Yes
Column Pruning	No	No	Yes	Yes	Yes
Predicate Pushdown	No	No	No	Yes	Yes
Indexing Information	No	No	No	Yes	Yes
Statistics Information	No	No	No	Yes	Yes
Nested Records	No	No	Yes	Yes	Yes



# Rule-based choice (heuristic)

Given a flow represented as  $\text{DAG}(V, E)$

- SequenceFile
  - $\text{size}(\text{getCol}(v)) = 2$
- Parquet
  - $\exists e \in O(v), \text{getType}(e) = \{\text{AggregationOps}\}$
  - $\exists e \in O(v), \text{getCol}(\text{getOP}(e)) \subset \text{getCol}(v)$
- Avro
  - $\forall e \in O(v), \text{getCol}(\text{getOP}(e)) = \text{getCol}(v)$
  - $\exists e \in O(v), \text{getType}(e) \in \{\text{Join}, \text{CartesianProduct}, \text{GroupAll}, \text{Distinct}\}$

# Cost-based choice

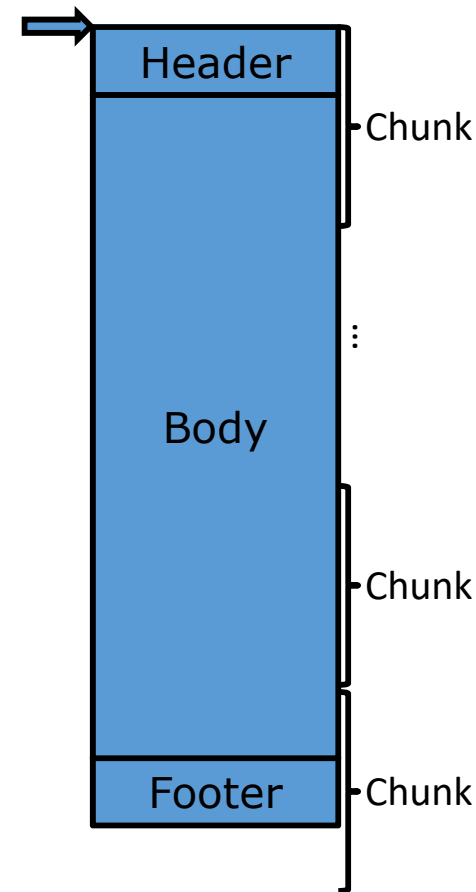
- Helps in choosing the right storage layout based on the workloads
- Costs to Consider
  - Write cost
  - Read Cost
    - Scan Operation
    - Projection Operation
    - Selection Operation
- Costs ignored
  - Block compression
  - Dictionary encoding (in Parquet)

# General formulas for size estimation

$$\begin{aligned} \text{Size}(\text{Layout}) = & \text{Size}(\text{Header}_{\text{Layout}}) \\ & + \text{Size}(\text{Body}_{\text{Layout}}) \\ & + \text{Size}(\text{Footer}_{\text{Layout}}) \end{aligned}$$

$$\text{UsedChunks}(\text{Layout}) = \frac{\text{Size}(\text{Layout})}{\text{Size}(\text{chunk})}$$

$$\text{Seeks}(\text{Layout}) = \lceil \text{UsedChunks}(\text{Layout}) \rceil$$



“Layout” can be either “Horizontal”, “Vertical” or “Hybrid”

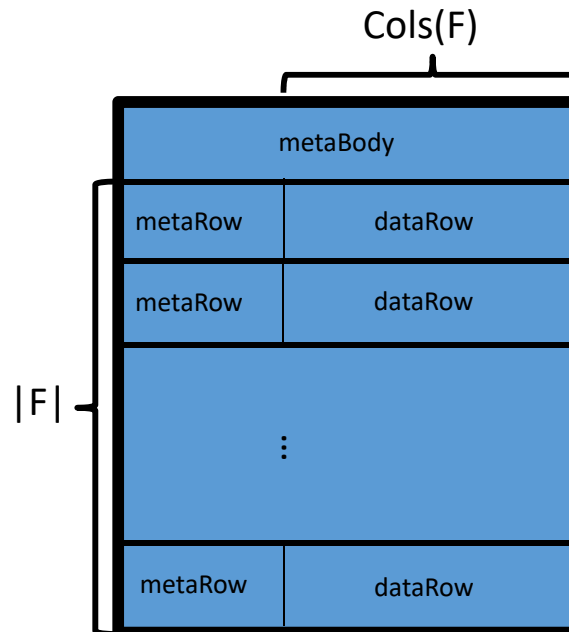


# Horizontal layout size estimation

$$\text{Size}(\text{dataRow}) = \text{Cols}(F) * \text{Size}(\text{cell})$$

$|F|$  - Number of rows  
 $\text{Cols}(F)$  - Number of columns  
 $\text{Size}(\text{cell})$  can be computed as average

$$\text{Size}(\text{Body}_{\text{Horizontal}}) = \text{Size}(\text{metaBody}) + |F| * (\text{Size}(\text{metaRow}) + \text{Size}(\text{dataRow}))$$



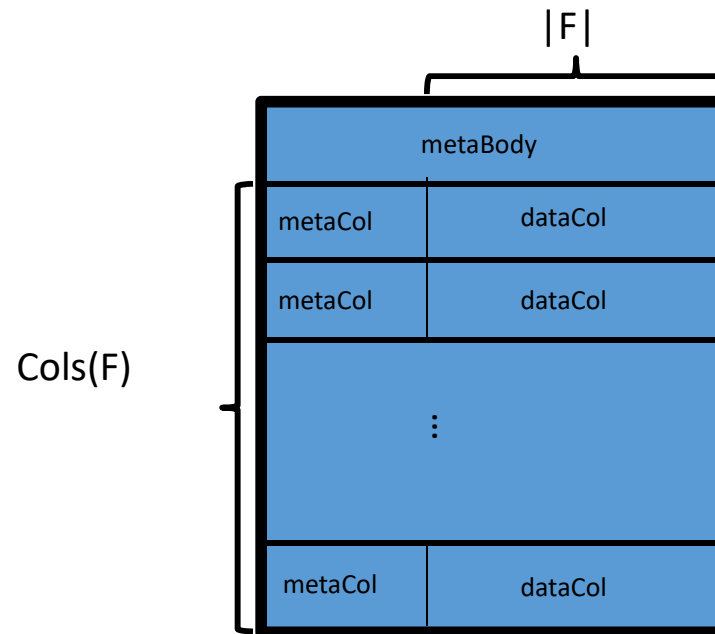
“F” is the content of the file

# Vertical layout size estimation

$$Size(dataCol) = |F| * Size(cell)$$

Size(cell) - Average cell (i.e., value) size

$$Size(Body_{Vertical}) = Size(metaBody) + Cols(F) * (Size(metaCol) + Size(dataCol))$$



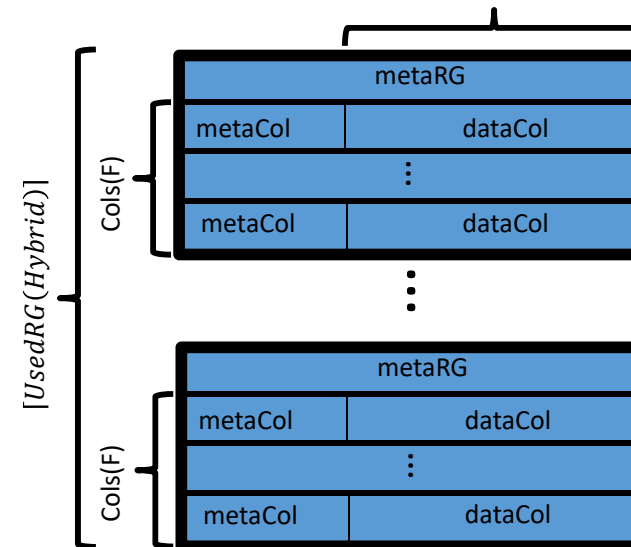
"F" is the content of the file

# Hybrid layout size estimation

$$UsedRG(Hybrid) = \frac{Cols(F) * |F| * Size(cell)}{(Size(RG) - Size(metaRG) - Cols(F) * Size(metaCol))}$$

$$Size(Body_{Hybrid}) = [UsedRG(Hybrid)] * (Size(metaRG) + Cols(F) * Size(metaCol)) + Cols(F) * |F| * Size(cell)$$

$$UsedRows(RG) = \frac{|F|}{UsedRG(Hybrid)}$$



“F” is the content of the file

# General formulas for cost estimation

$W_{writeTransfer}$  - Weight of transferring data

$1 - W_{writeTransfer}$  - Weight of positioning arm

If weight of transferring data is as expensive as positioning arm, both value will be 0.5

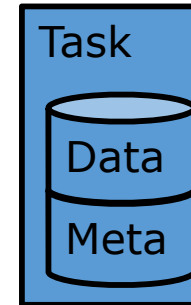
If weight of transferring data is as double expensive as positioning arm, one will be 0.66 another 0.33

$$Cost(Write_{Layout}) = UsedChunks(Layout) * W_{WriteTransfer} + Seeks(Layout) * (1 - W_{WriteTransfer})$$

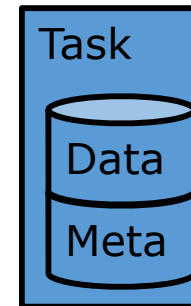
$$Size(Scan_{Layout}) = Size(Layout) + ([UsedChunks(Layout)] * Size(Meta_{Layout}))$$

$$UsedChunks(Scan_{Layout}) = \frac{Size(Scan_{Layout})}{Size(chunk)}$$

$$Cost(Scan_{Layout}) = UsedChunks(Scan_{Layout}) * W_{ReadTransfer} + Seeks(Layout) * (1 - W_{ReadTransfer})$$



...



"Layout" can be either "Horizontal", "Vertical" or "Hybrid"

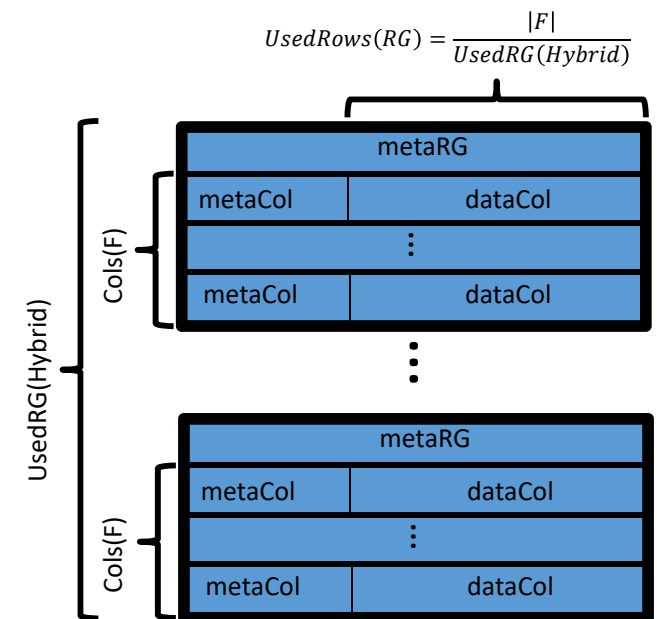
# Projection in hybrid layouts

This is only possible with hybrid layout  
With horizontal layout, we only read whole data

$$Size(projCols) = Proj(F) * UsedRows(RG) * Size(cell)$$

$$Size(Project_{Hybrid}) = Size(Header_{Hybrid}) + Size(Footer_{Hybrid}) \\ + [UsedRG(Hybrid)] * (Size(metaRG) + proj(F) * Size(metaCol)) \\ + UsedRG(Hybrid) * Size(projCols)$$

$$Cost(Project_{Hybrid}) = UsedChunks(Project_{Hybrid}) * W_{ReadTransfer} \\ + Seeks(Hybrid) * (1 - W_{ReadTransfer})$$



“Proj(F)” is the number of projected columns

# Probability of retrieving a RowGroup

- Probability of a row fulfilling P (a.k.a. selectivity factor)  
SF
- Probability of a row NOT fulfilling P  
 $1 - SF$
- Probability of none of the rows in a RowGroup fulfilling P  
 $(1 - SF) \cdot (1 - SF) \cdot \dots \cdot (1 - SF) = (1 - SF)^{\text{UsedRows}(\text{RG})}$
- Probability of some row in a RowGroup fulfilling P  
 $1 - (1 - SF)^{\text{UsedRows}(\text{RG})}$

# Selection in hybrid layouts

$$P(RGSelected) = 1 - (1 - SF)^{UsedRows(RG)}$$

$$Size(Select_{Hybrid}) = \left[ \frac{SF * |F|}{UsedRows(RG)} \right] (Size(metaRG) + Cols(F) * Size(metaCol)) \quad \text{metadata}$$

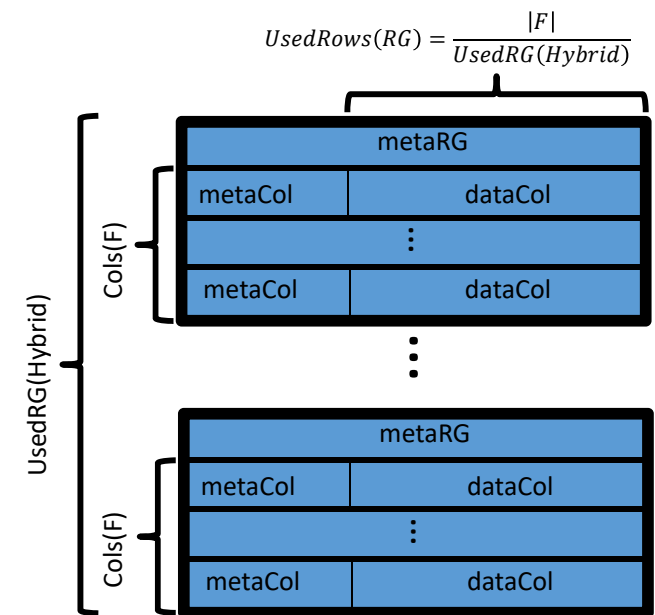
$$+ \underbrace{SF * |F| * Cols(F) * Size(cell)}_{\text{data}}$$

$$UsedRG(Select_{Hybrid}) = \begin{cases} \text{if unsorted: } P(RGSelected) * UsedRG(Hybrid) \\ \text{if sorted: } \left\lfloor \frac{Size(Select_{Hybrid})}{Size(RG)} \right\rfloor \end{cases}$$

$$Size(Select_{Hybrid}) = Size(Header_{Hybrid}) + Size(Footer_{Hybrid}) + UsedRG(Select_{Hybrid}) * Size(RG)$$

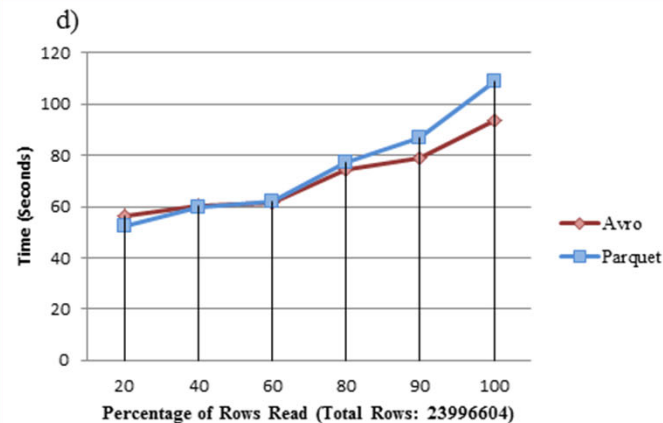
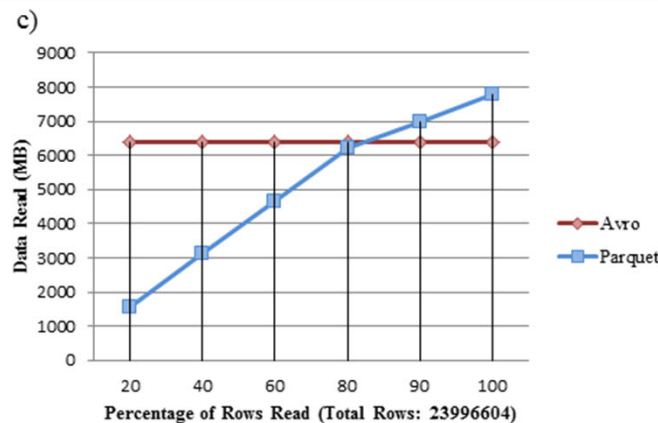
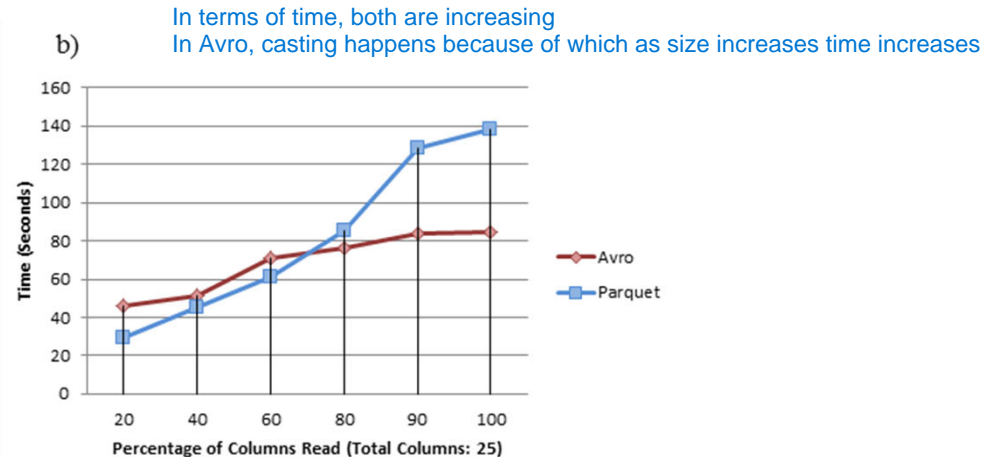
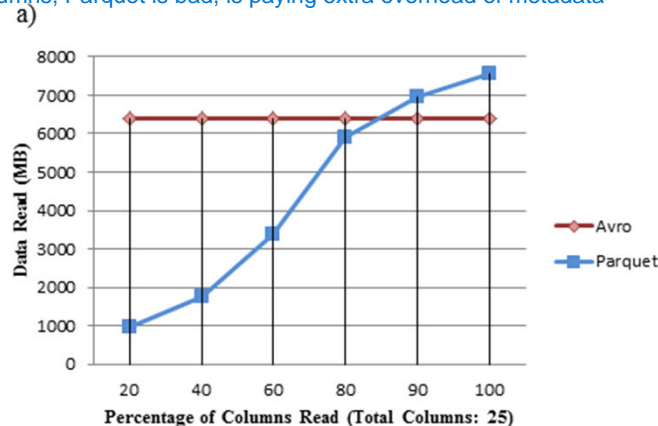
$$Cost(Select_{Hybrid}) = UsedChunks(Select_{Hybrid}) * W_{ReadTransfer} + Seeks(Select_{Hybrid}) * (1 - W_{ReadTransfer})$$

“F” is the content of the file



# Comparison of selection and projection

No matter, how much columns we are retrieving, for Avro it is always same  
 If we are retrieving few columns, Parquet is good  
 If we are retrieving all columns, Parquet is bad, is paying extra overhead of metadata

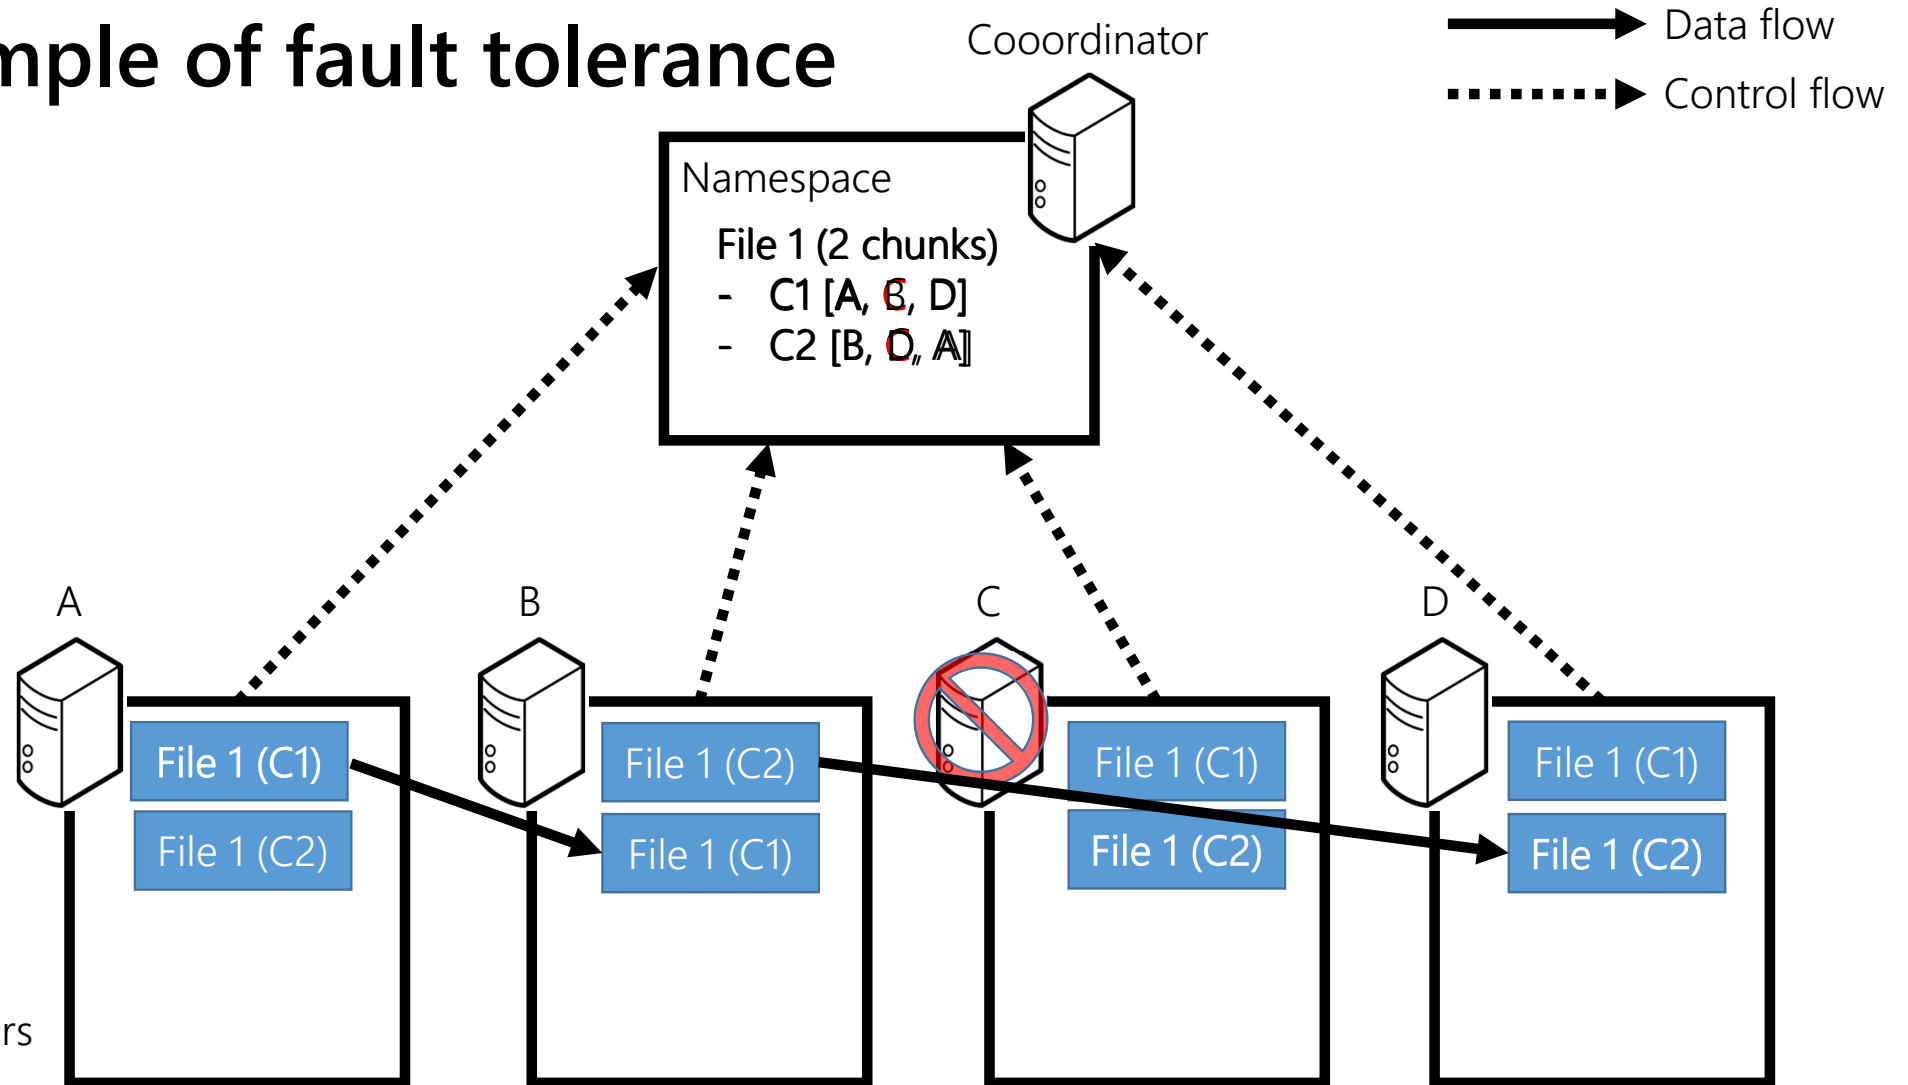




# Fault tolerance

- Managed from the coordinator
    - It expects to receive every 3 seconds a *heartbeat* message from chunkservers
  - Chunkserver not sending a heartbeat for 60 seconds, a fault is declared
  - Corrective actions When chunk server crashes, two things happen
    - Update the namespace
    - Copy one of the replicas to a new chunkserver Everything that was available in this chunk server is recreated in another chunk server to maintain the number of replicas created
      - Potentially electing a new primary replica
- If there was some primary replica, it will have some sort of election to find new primary replica

# Example of fault tolerance



# (Distributed) Catalog Management

Challenge II

# Client caching

## Cache miss

1. The client sends a READ command to the coordinator
2. The coordinator requests chunkservers to send the chunks to the client
  - Ranked according to the closeness in the network
3. The list of locations is cached in the client
  - Not a complete view of all chunks

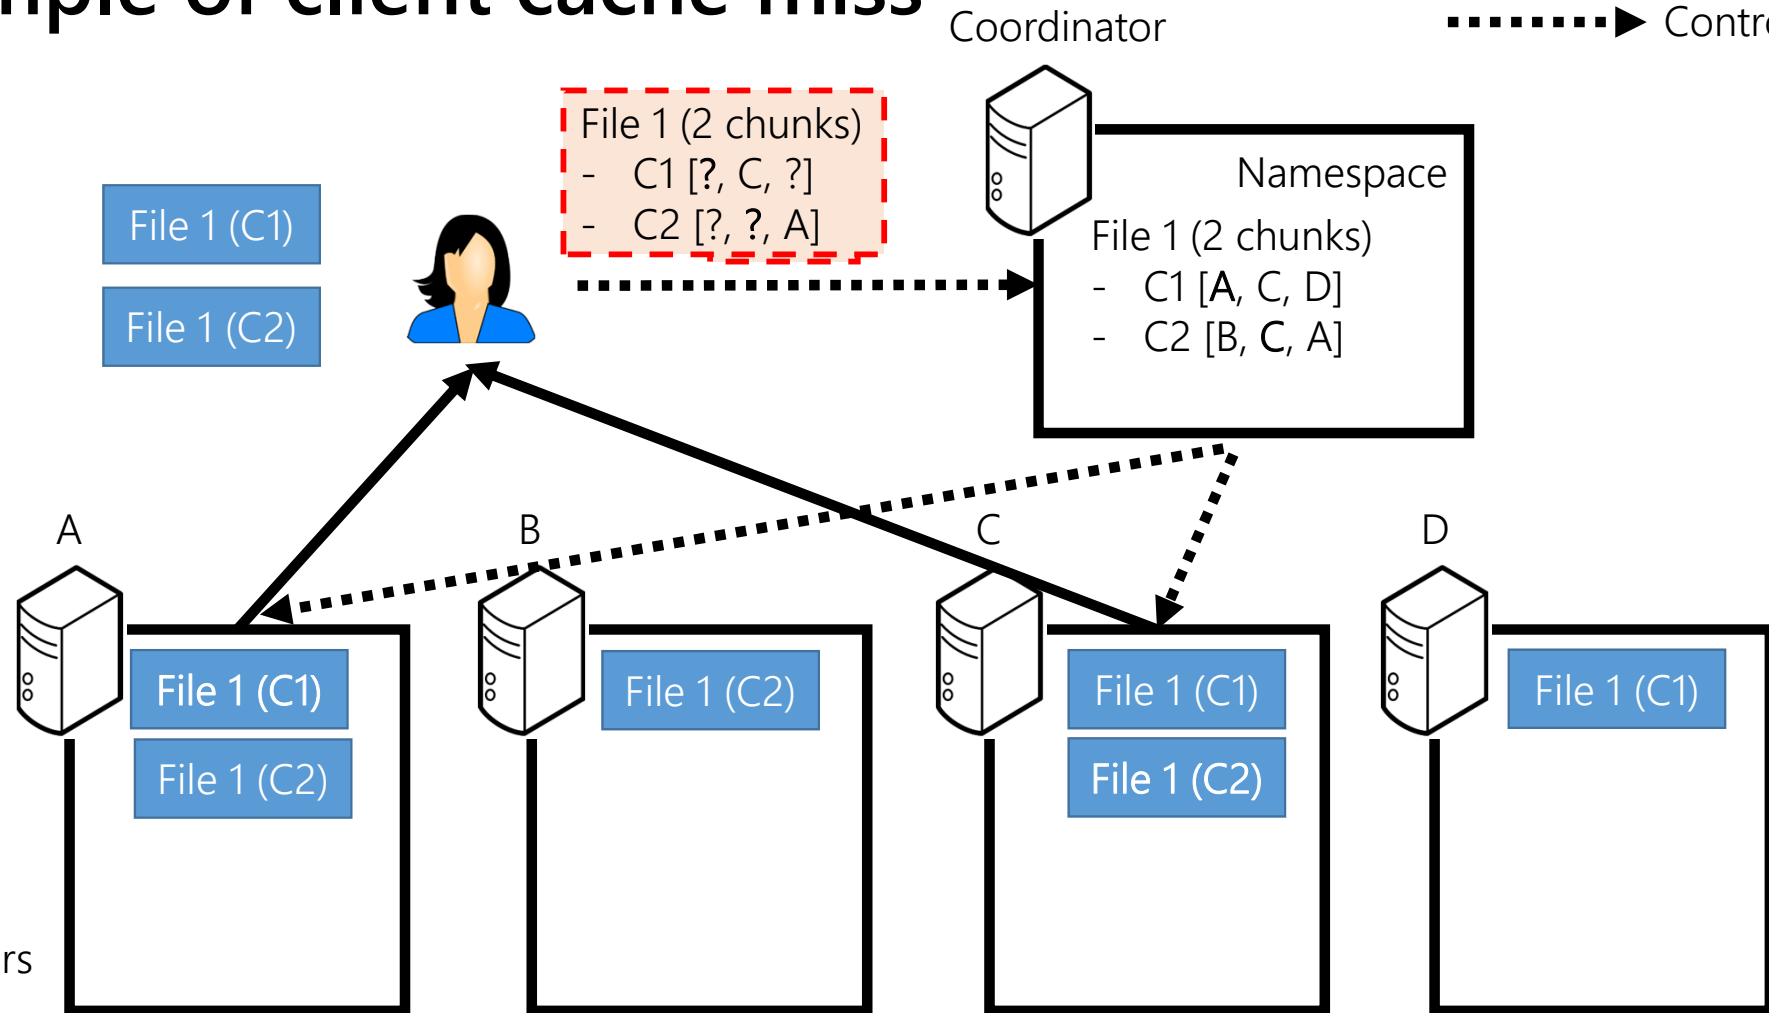
## Cache hit

1. The client reads the cache and requests the chunkservers to send the chunks

Avoid coordinator bottleneck  
+  
One communication step is saved

# Example of client cache miss

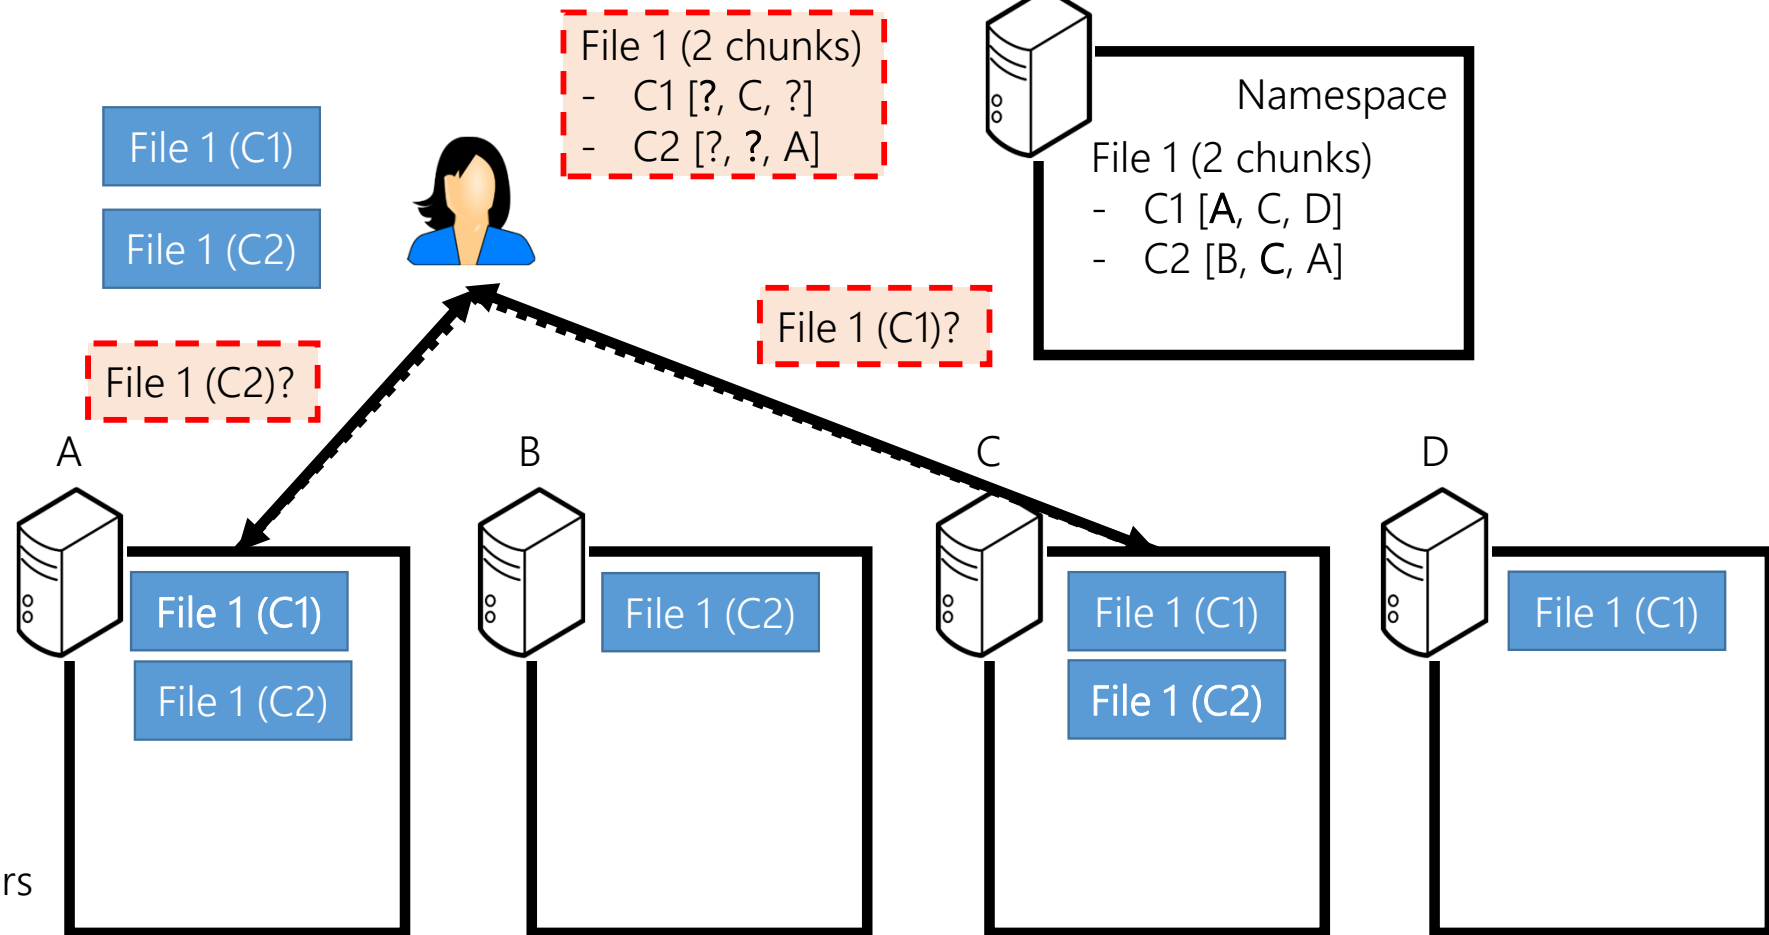
—————▶ Data flow  
.....▶ Control flow



ChunkServers

# Example of client cache hit

—————▶ Data flow  
.....▶ Control flow



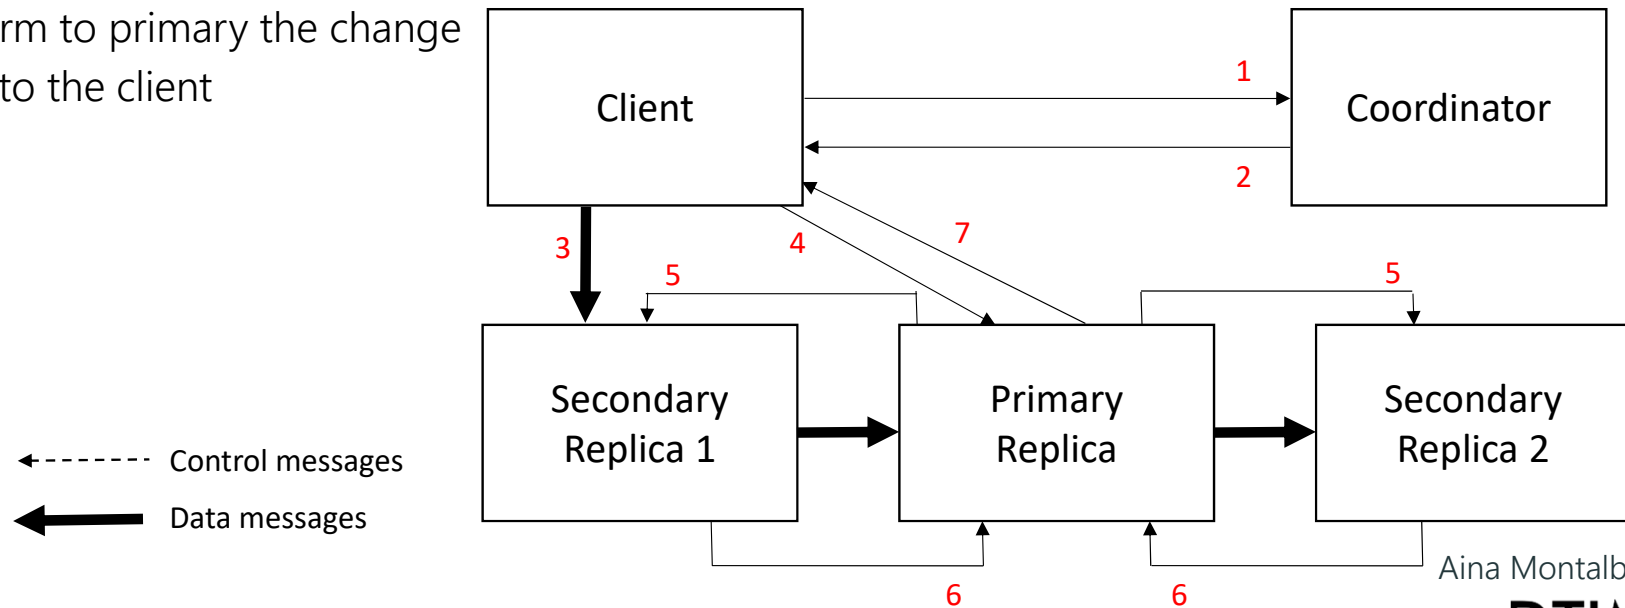
ChunkServers

# (Distributed) Transaction Management

Challenge III

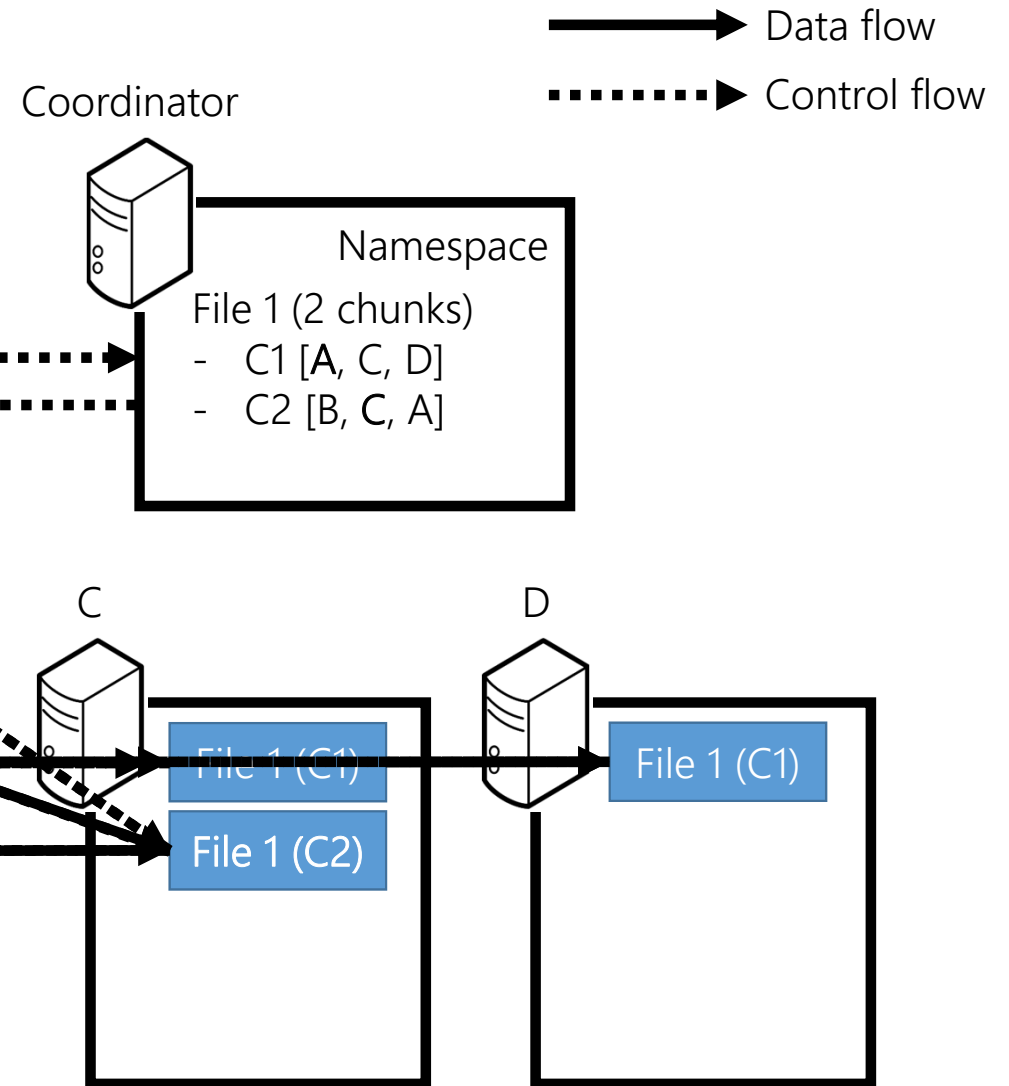
# Writing replicas

1. The client requests the list of the replicas of a file
2. Coordinator returns metadata
3. Client sends a chunk to the closest chunkserver in the network
  - This chunk is pipelined to the other chunkservers in the order defined by the master (leases)
4. Client sends WRITE command to primary replica
5. Primary replica sends WRITE command to secondary replicas
6. Secondaries confirm to primary the change
7. Primary confirms to the client





# Example of writing replicas

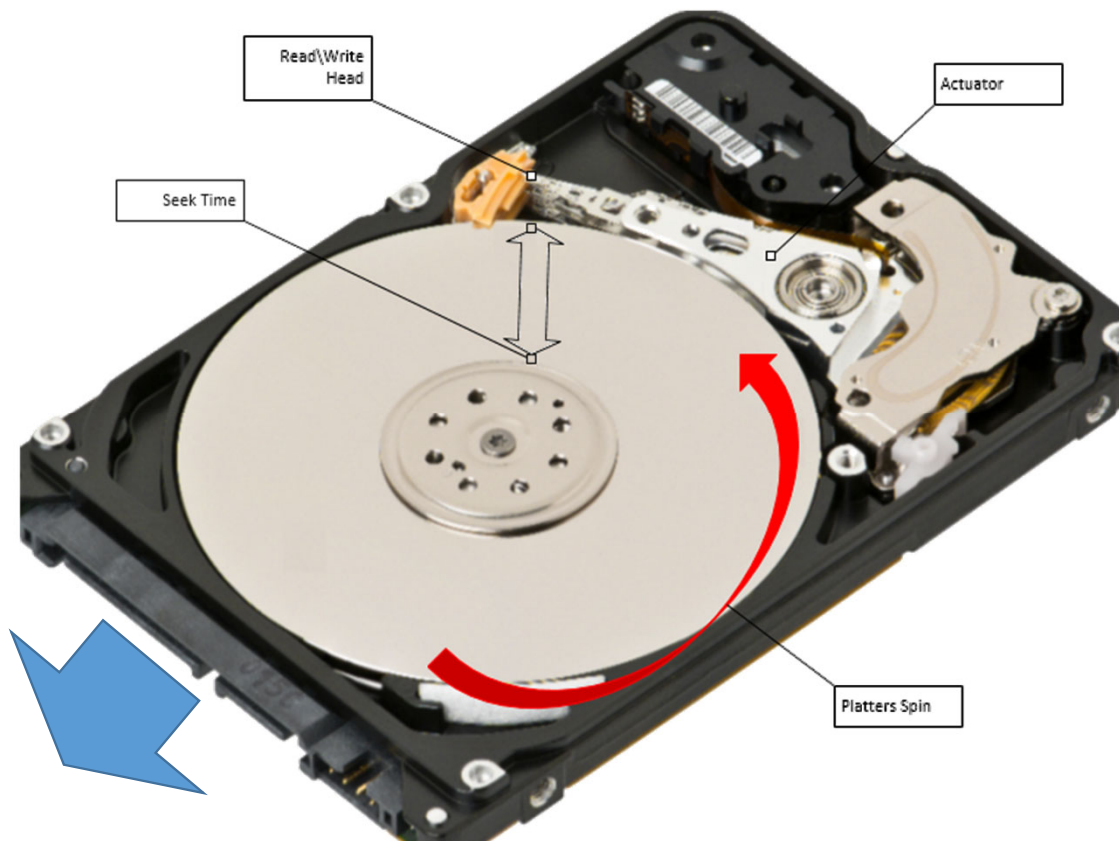


ChunkServers

# (Distributed) Query processing

Challenge IV

# HDDs costs



## Rotational Latency

The amount of time taken for the platters to spin the data under the head (measured in RPM)

## Seek Time

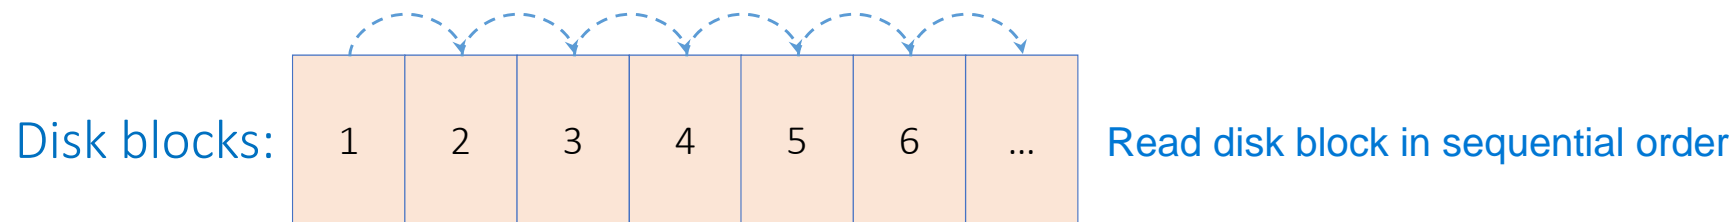
Time taken for the ReadWrite head (mechanical arm) to move between cylinders on the disk

## Transfer Time

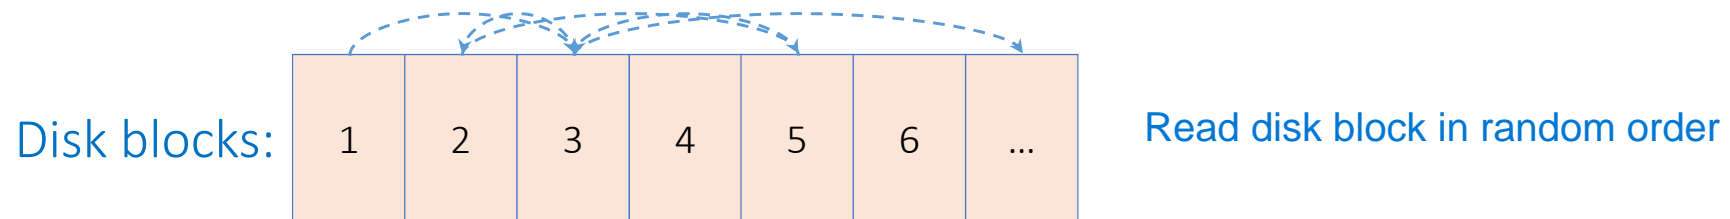
Time taken for requests to get from the system to the disk (depends on the block size, e.g., 8KB)

# Sequential vs. Random access

## Sequential Access



## Random Access



Transfer time is same in both type of access  
But seek time and rotational latency might be different

# Cost of accessing data (approximations)

$n$  - number of blocks to be fetched in memory

- Sequential reads

- Option to maximize the effective read ratio
  - Depends on DB design
- Enables pre-fetching enables use of prefetching mechanism  
knows which block to fetch next

Cost = seek + rotation +  $n \times$  transfer

- Random Access

- Requires indexing structures we cannot predict next block  
so prefetching is not possible  
To know which is the next block we are going to read
- Ignores data locality After reading one block, we are not going to read immediate next one

Cost<sub>single cylinder files</sub> = seek +  $n \times$  (rotation + transfer)

Cost<sub>multi-cylinder files</sub> =  $n \times$  (seek + rotation + transfer)

If all data is in same cylinder,  
1 seek

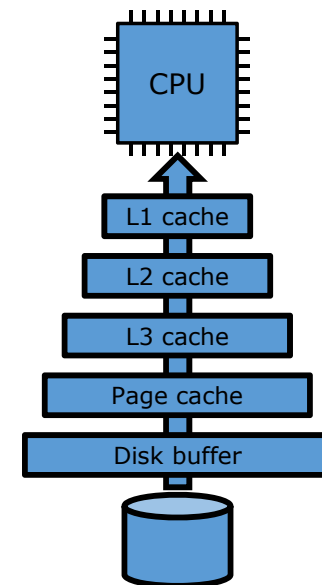
For each block, wait for rotation and transfer data

If data is in different cylinder  
For each block, perform seek

wait for rotation and transfer the data

Seek is most expensive

seek	~12ms
rotation	~3ms
transfer (8KB)	~0.03ms



# Closing

# Summary

- GFS architecture and components
- GFS main operations
  - Fault tolerance
  - Writing files and maintenance of replicas
  - Reading files
- HDFS file formats
  - Horizontal
  - Vertical
  - Hybrid

# References

- S. Ghemawat et al. *The Google File System*. OSDI'03
- K. V. Shvachko. *HDFS scalability: the limits to growth*. 2010
- S. Abiteboul et al. *Web data management*. Cambridge University Press, 2011
- A. Jindal et al. *Trojan data layouts: right shoes for a running elephant*. SOCC, 2011
- F. Färber et al. *SAP HANA database: data management for modern business applications*. SIGMOD, 2011
- V. Raman et al. *DB2 with BLU Acceleration: So Much More than Just a Column Store*. VLDB, 2013
- D. Abadi, et al. *Column-stores vs. row-stores: how different are they really?* SIGMOD Conference, 2008
- M. Stonebraker et al. *C-Store: A Column-oriented DBMS*. VLDB, 2005
- G. Copeland and S. Khoshafian. *A Decomposition Storage Model*. SIGMOD Conference, 1985
- F. Munir. *Storage Format Selection and Optimization of Materialized Intermediate Results in Data-Intensive Flows*. PhD Thesis (UPC), 2019
- A. Hogan. *Procesado de Datos Masivos* (U. Chile)