

# Big Data Architectures

Big Data Management

# Knowledge objectives

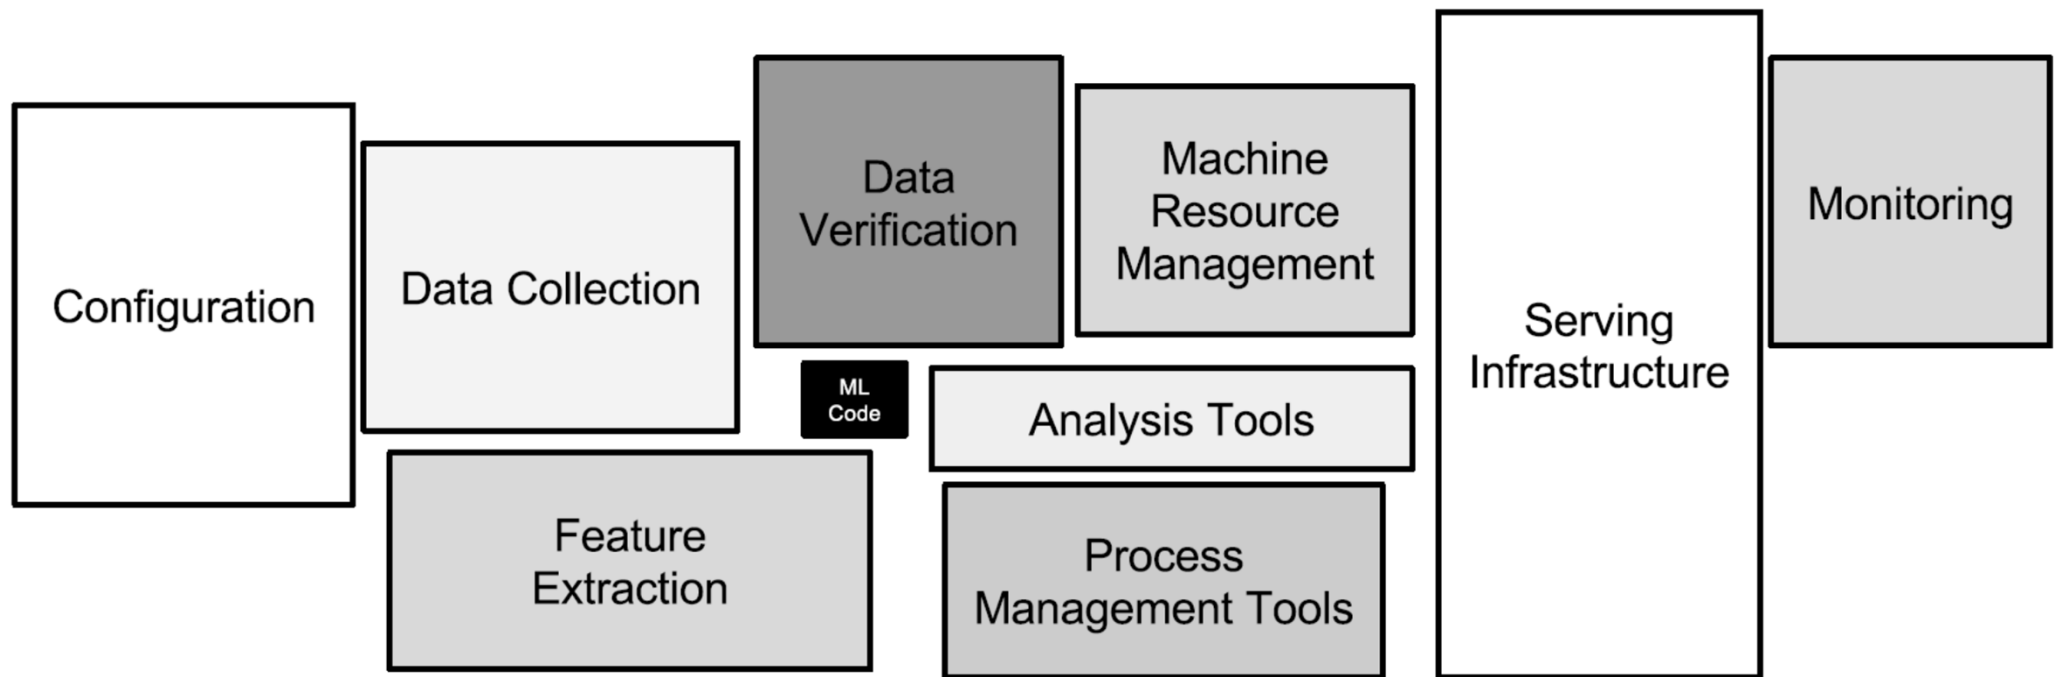
1. Explain the problema of a spaghetti architecture
2. Explain the need of the Lambda architecture
3. Explain the difference between the Kappa and Lambda architectures
4. Justify the need of a Data Lake
5. Identify the difficulties of a Data Lake
6. Explain the need of each component in the Bolster architecture
7. Map the components of Bolster to a RDBMS architecture

# Application Objectives

1. Given a use case, define its software architecture

# Problem definition

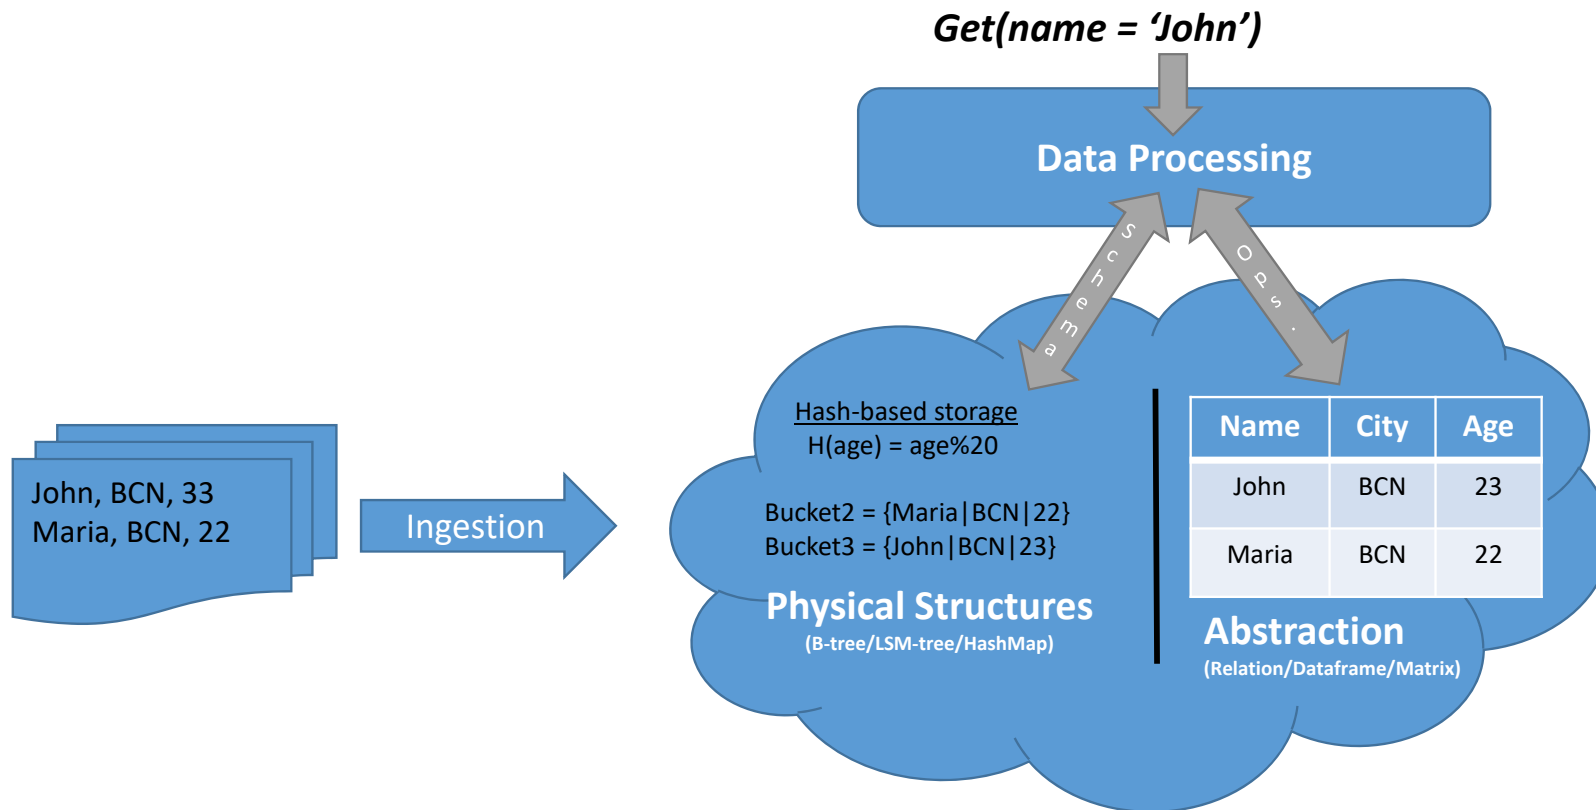
# Vast and Complex surrounding infrastructure



- ML code is just a small box in the middle
  - ... with a lot of plumbing around it

D. Sculley et al.

# Data Management (I)



# Data Management (II)

**Data management** refers to the functionalities a DBMS must provide:

- **Ingestion**: means provided to insert /upload data
  - E.g., ORACLE SQL\*Loader
- **Storage**: format/structures used to persist data
  - E.g., hash, B-tree, heap file
- **Modelling**: arrangement of data within the available structures
  - E.g., normalization, partitioning
- **Processing**: means provided to manipulate data
  - E.g., PL/SQL
- **Querying/fetching**: means provided to allow users to retrieve data
  - E.g., SQL, Relational Algebra

In **Big Data settings**, they are the **same** concepts but assuming NOSQL underneath

1. Typically, a distributed system
2. Possibly with an alternative data model to the Relational one
3. Implementing ad-hoc architectural solutions

# Big Data Architectures

- Question the main principles of traditional DB architectures
  - Data can grow beyond limits requiring scale out (a.k.a. **Volume**)
  - Data is not necessarily persisted (a.k.a. **Velocity**)
  - Data structure is neither known a priori, nor fixed (a.k.a. **Variety** and Variability)
- Use new trendy technological features
  - Primary indexes to implement the global catalog
    - Distributed Tree
    - Dynamic Hashing
  - In-memory processing
  - Columnar block iteration: vertical fragmentation + fixed-size values + RLE compression
    - Heavily exploited by column-oriented databases
    - Good for read-only workloads
  - Sequential reads for large workloads
    - Take the most out of databases by boosting sequential reads
      - Enables pre-fetching
      - Option to maximize the effective read ratio (by a good DB design)
        - Key design
- Implement from scratch the whole stack
  - Ingestion, Storage, Modeling, Processing, and Querying



# The Multi-Project Approach

- The DBMS tasks can be spread over different systems
  - Independent
  - Heterogeneous
- Hadoop is a paradigmatic case:
  - Storage: HDFS + Hbase
  - Modeling: HCatalog
  - Ingestion: Sqoop
  - Processing: Spark
  - Querying: Spark SQL

## Orchestrator

1. **Oozie is simple**
2. **Musketeer is a bit complex**
3. **Its role would be similar to that of a global execution manager.**

Ingestion - Sqoop, Kafka, Flume

Storage - HDFS+HBASE, mongodb

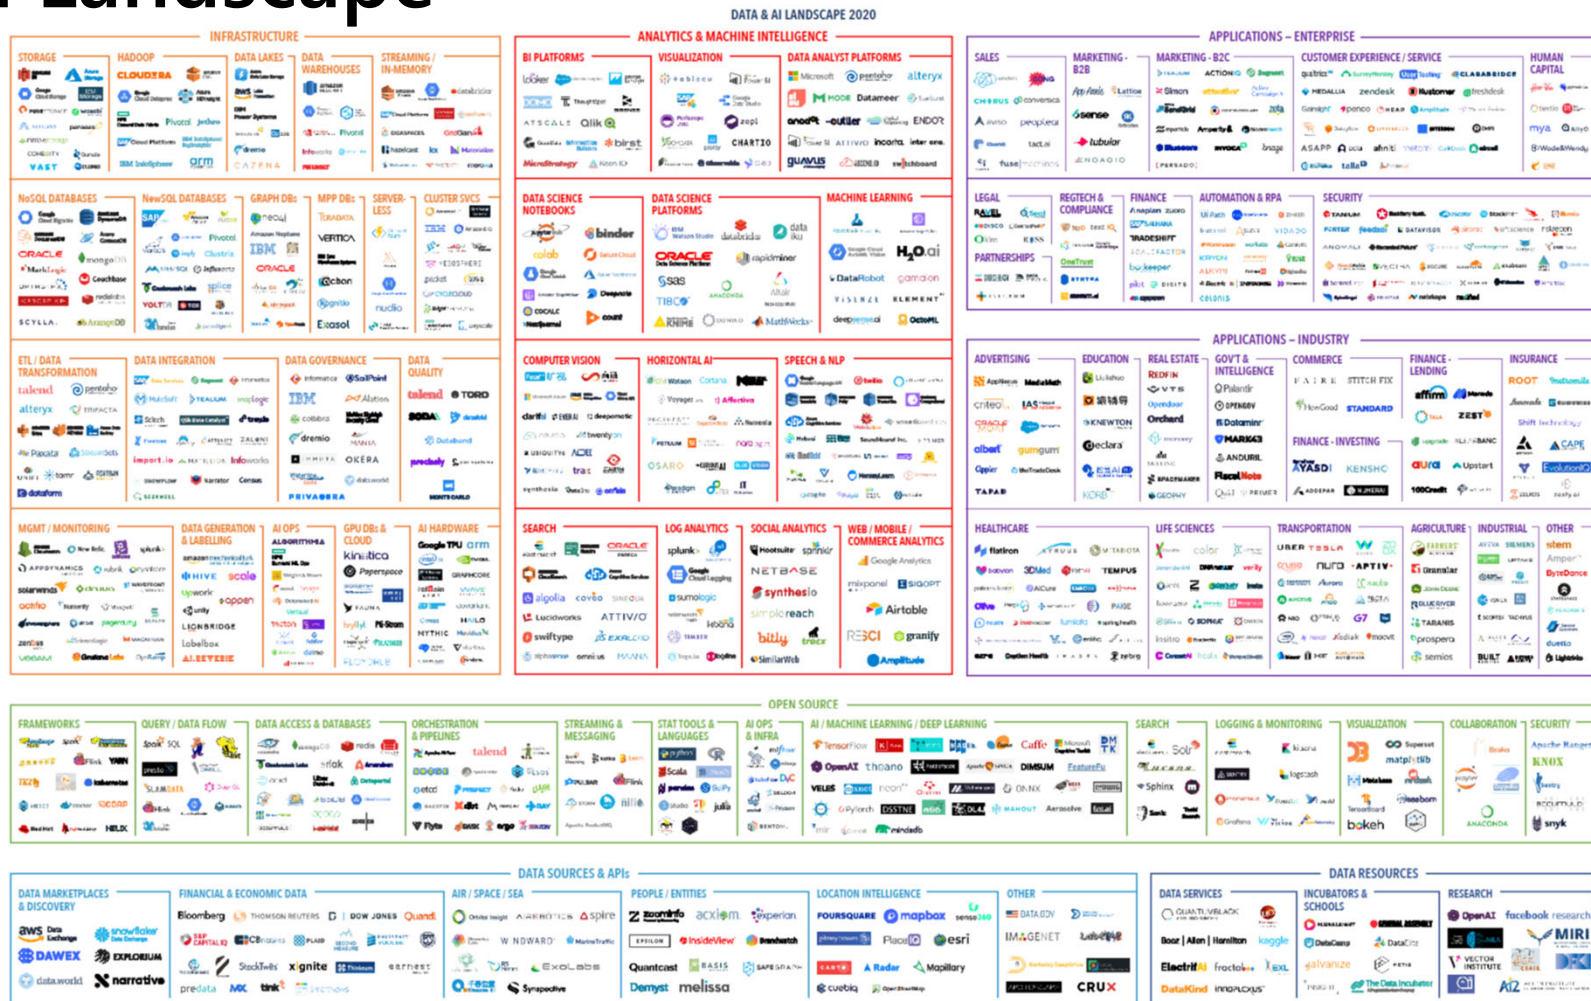
Processing - Spark, Flink, Spark Streaming, hadoop

metadata - hcatalog, ontop, graphdb, poolparty, jeena

Querying - Postgres, SparkSQL, Neo4j, redis, graphite

Analytics - QlikView, HIVE, SparkSQL, Big analytics, ml flow, mahout

# Big Data Landscape

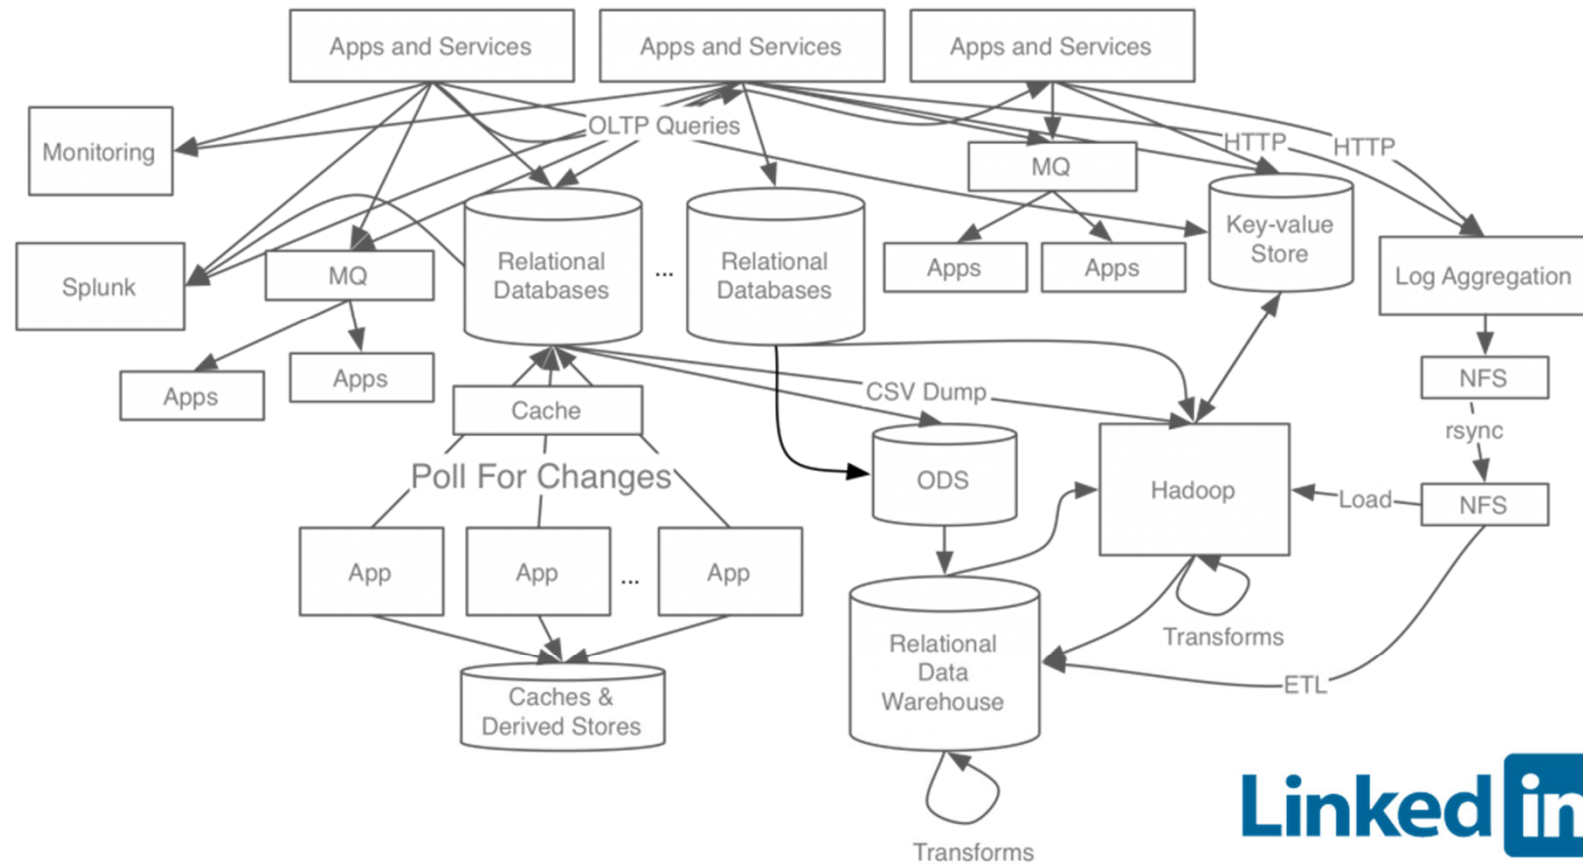


Version 1.0 - September 2020

© Matt Turck (@mattturck) & FirstMark (@firstmarkcap)

mattturck.com/data2020

# Spaghetti architecture



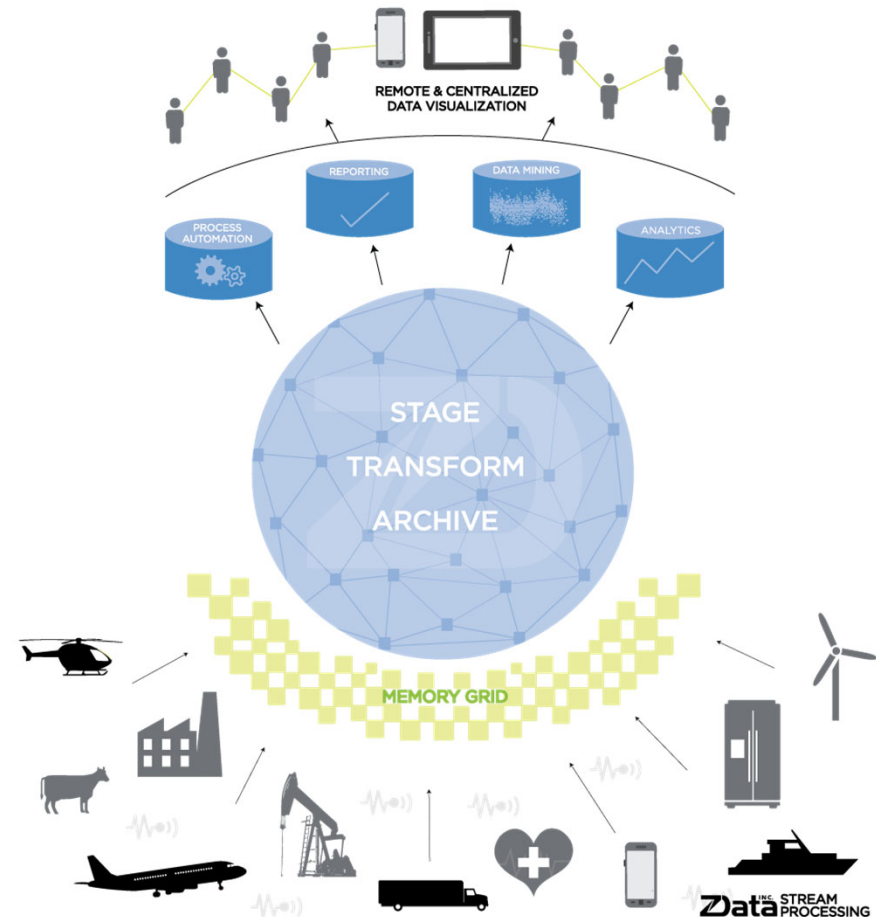
<https://www.confluent.io/blog/event-streaming-platform-1>

# New Storage Architectural Pattern

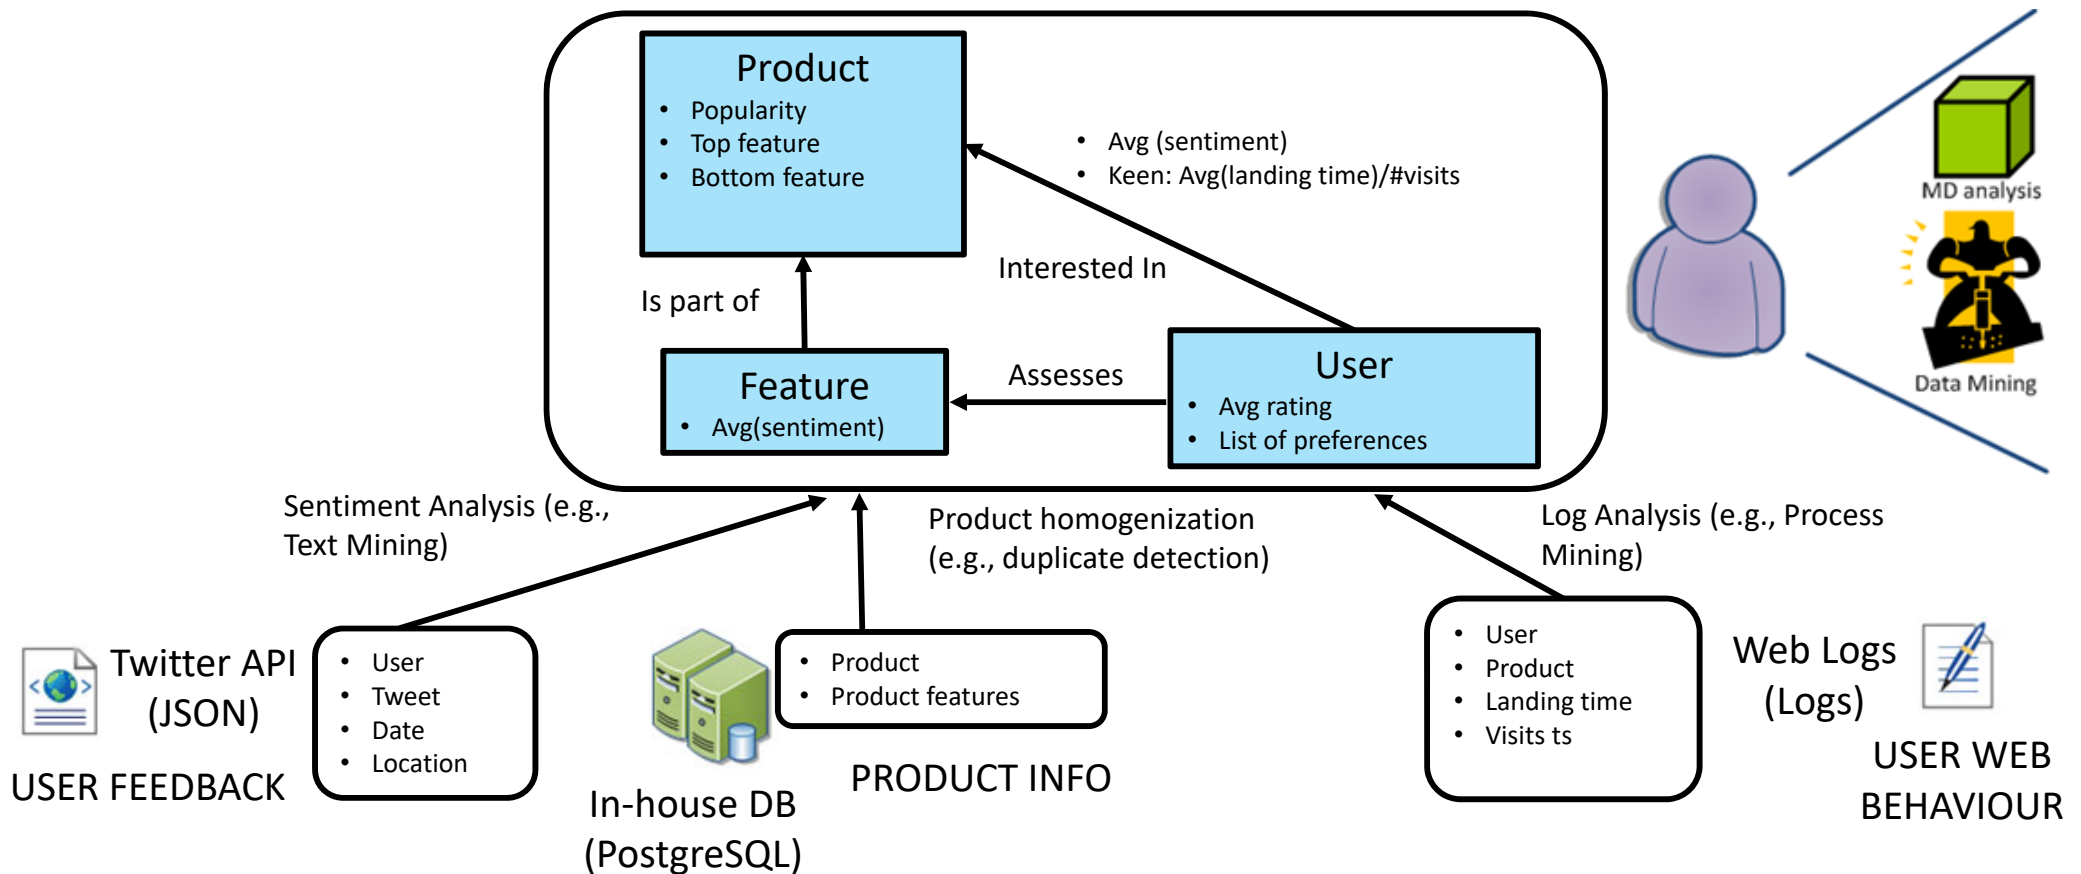
From data warehousing to data lakes

# The Data Lake

- Idea: Load-First, Model-Later
- Modelling at load time restricts the potential analysis that can be done later (Big Analytics)
- Characteristics:
  - a) Store raw data
  - b) Provide governing functionalities
  - c) Create on-demand views to handle precise analysis needs




# Model-First (Load-Later)

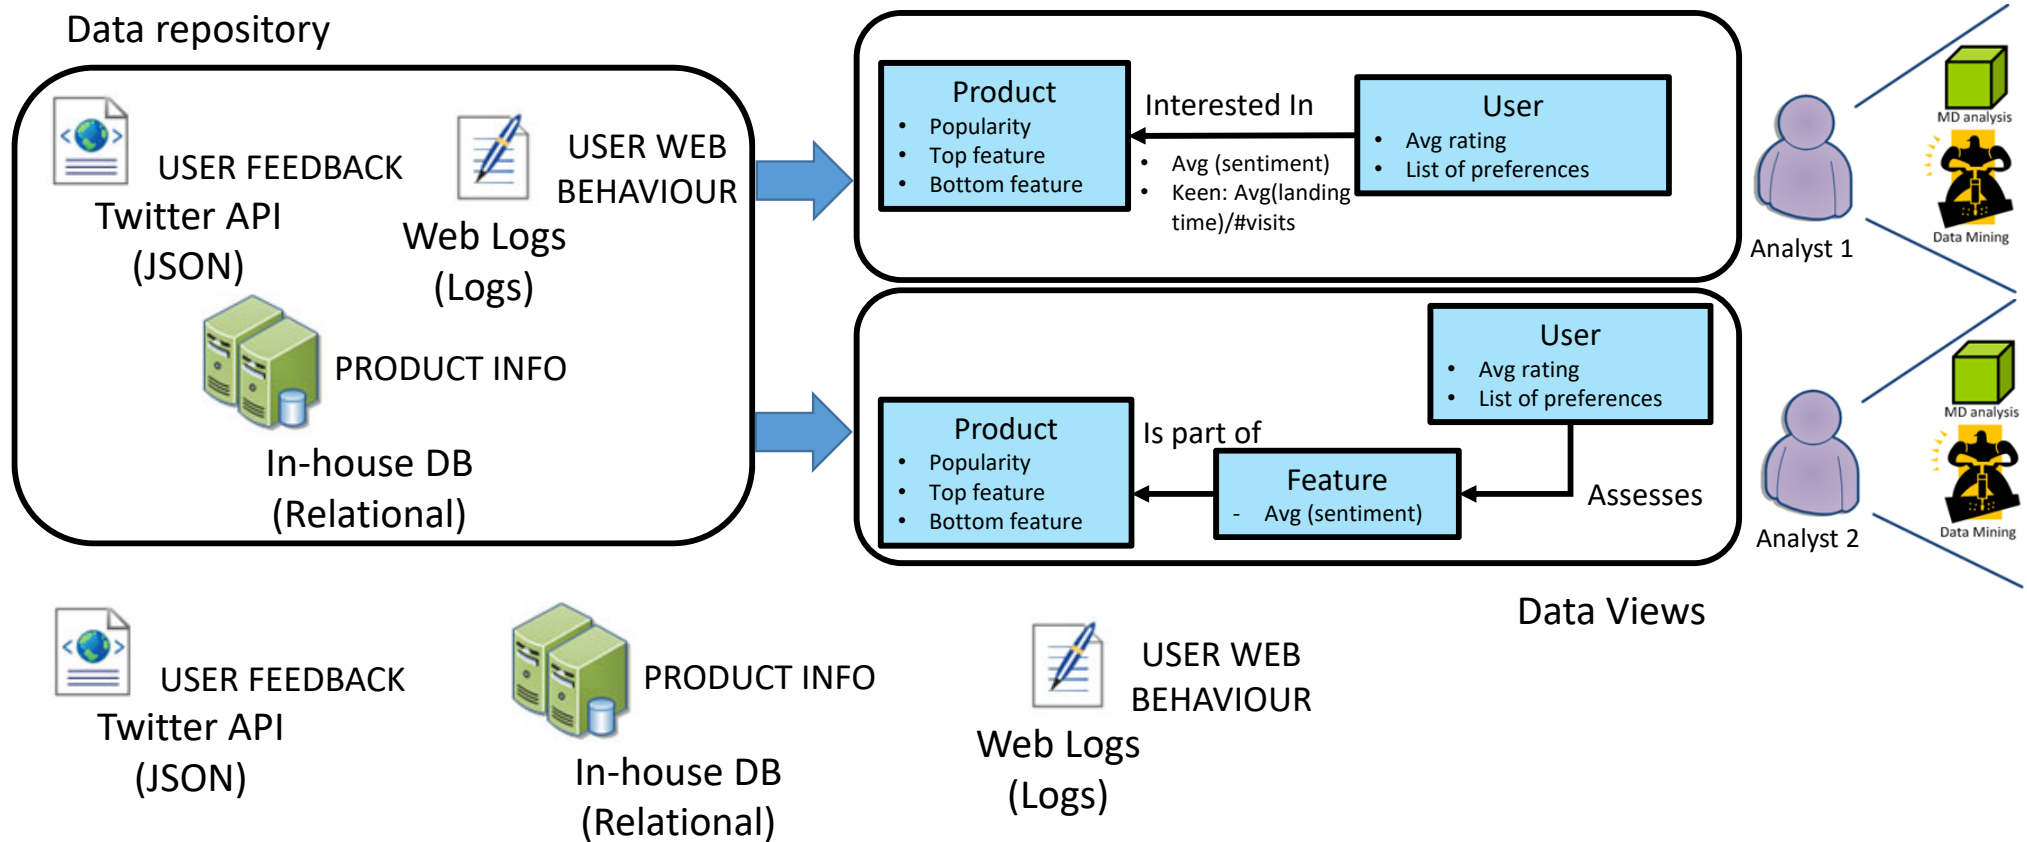




# Drawbacks of Model-First (Load-Later)



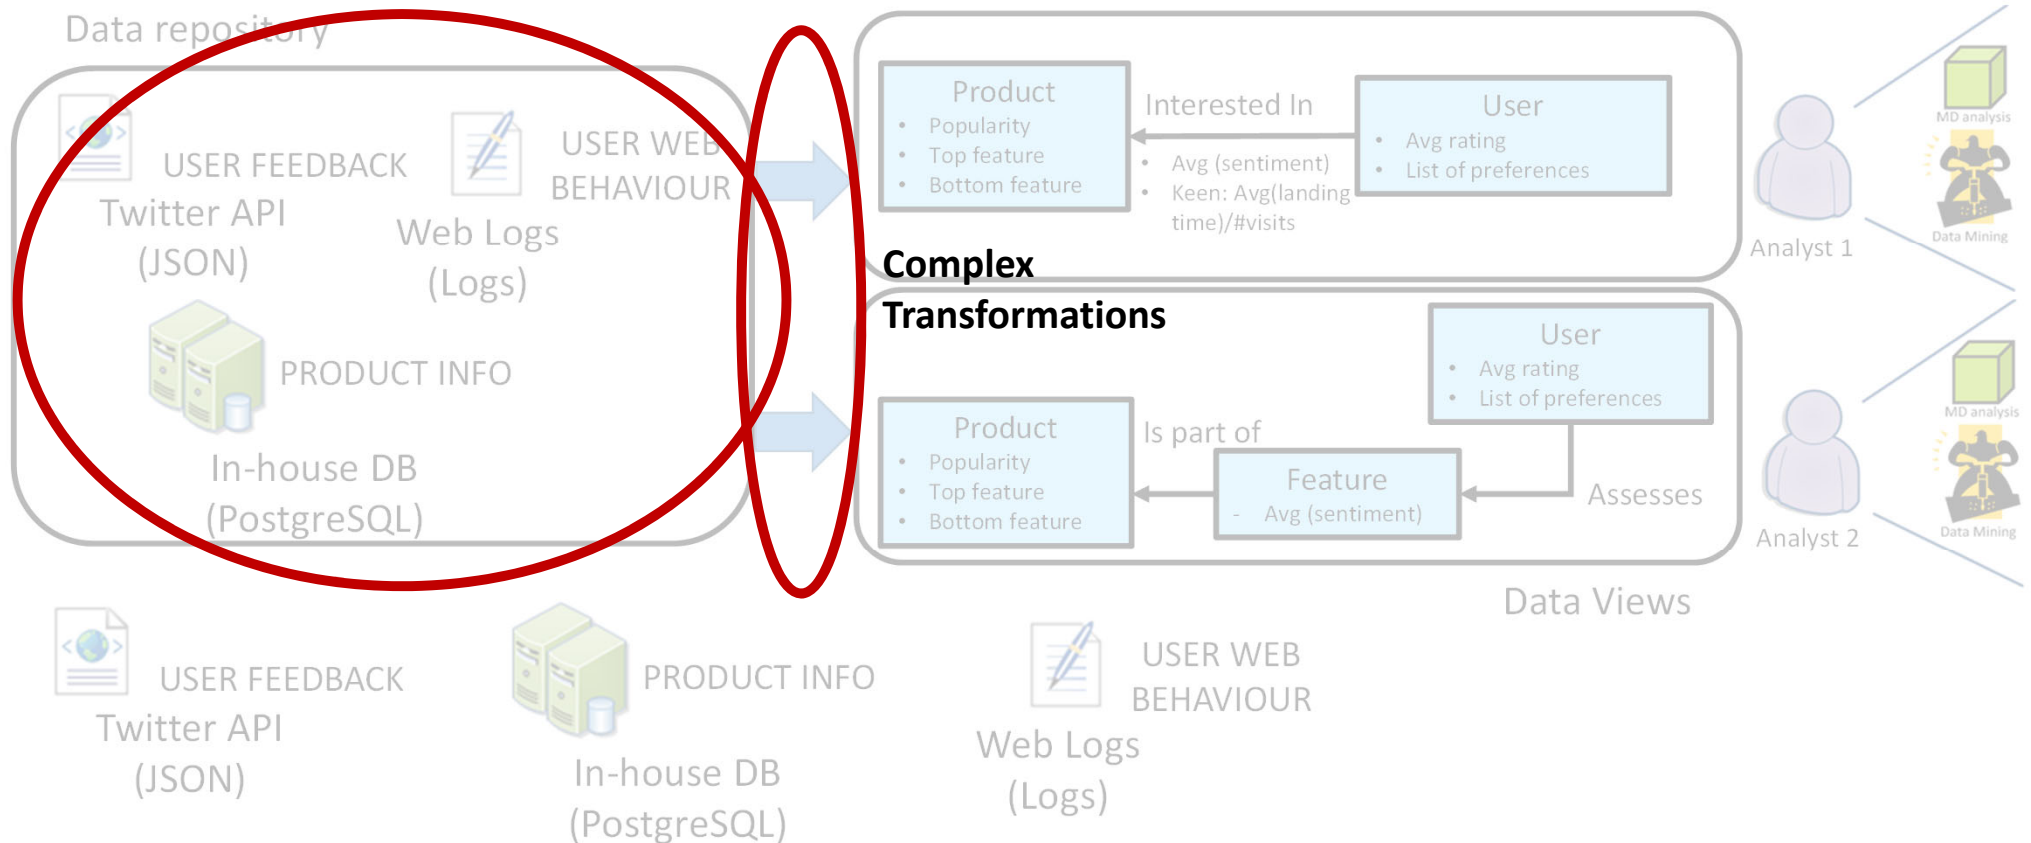
# Load-First (Model-Later)





# Drawbacks of Load-First (Model-Later)

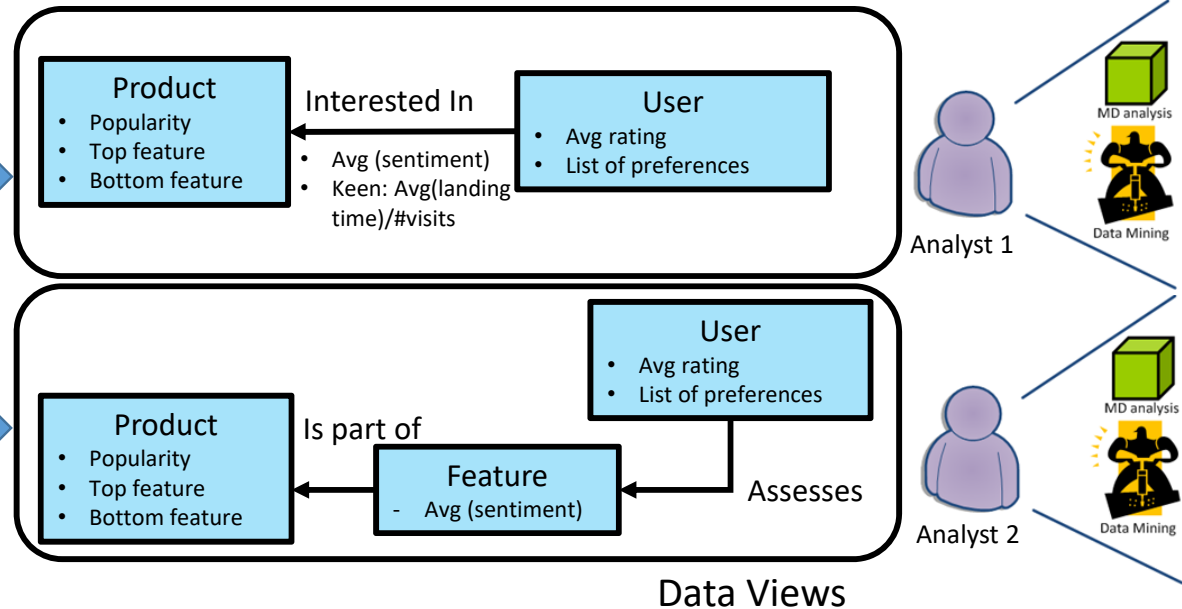
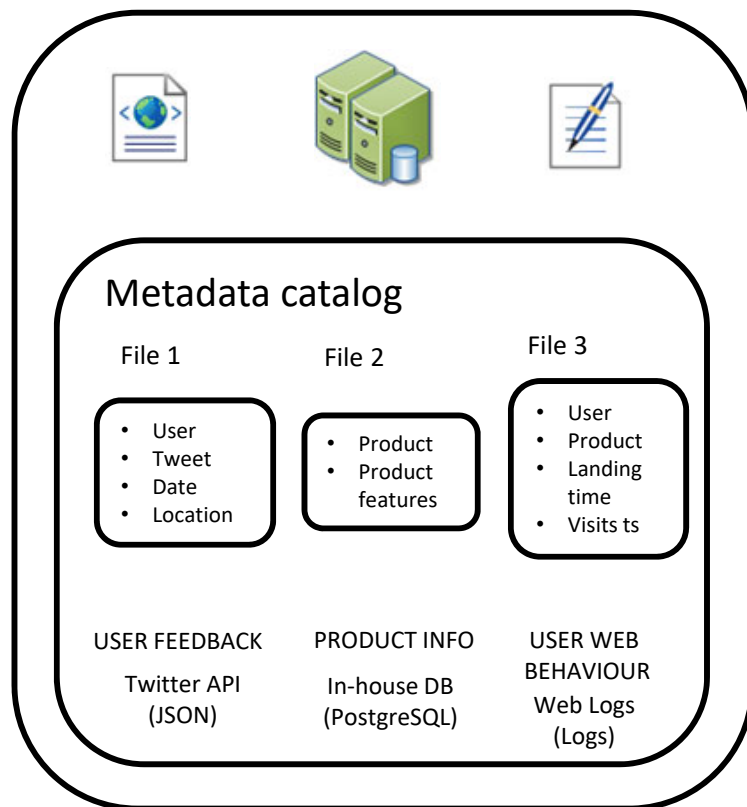
## Data Swamp



Stonebraker (2014)

# Towards semantic-awareness

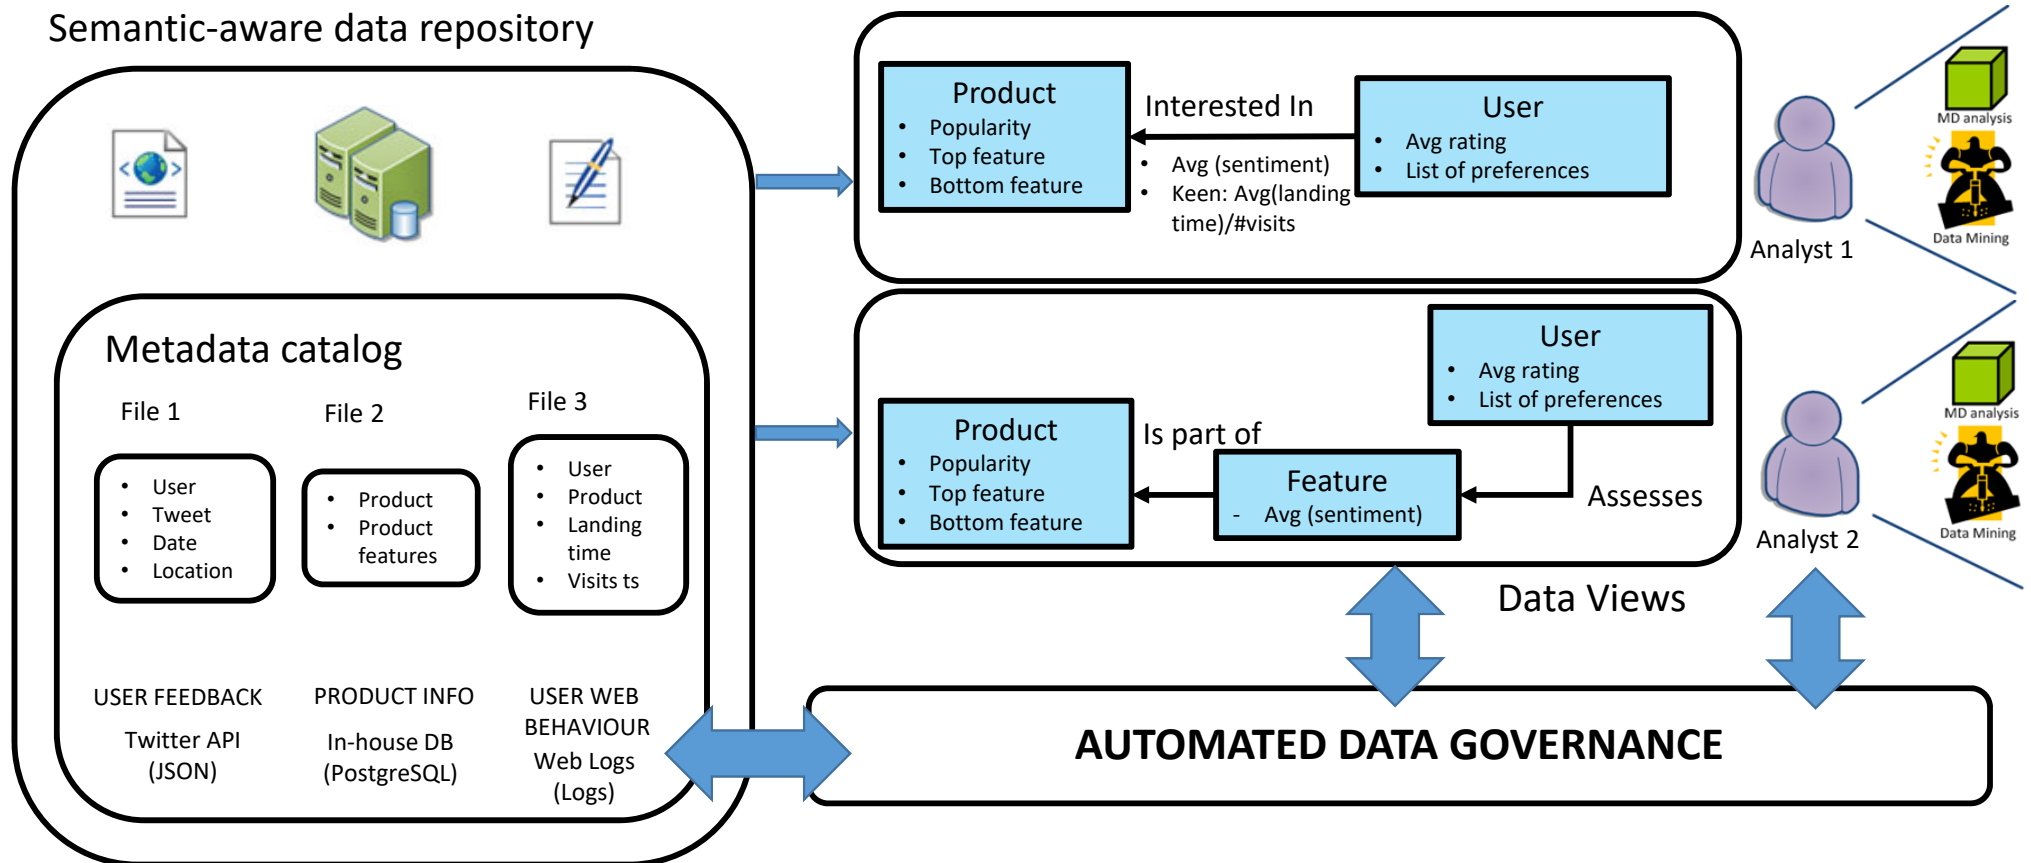
Semantic-aware data repository



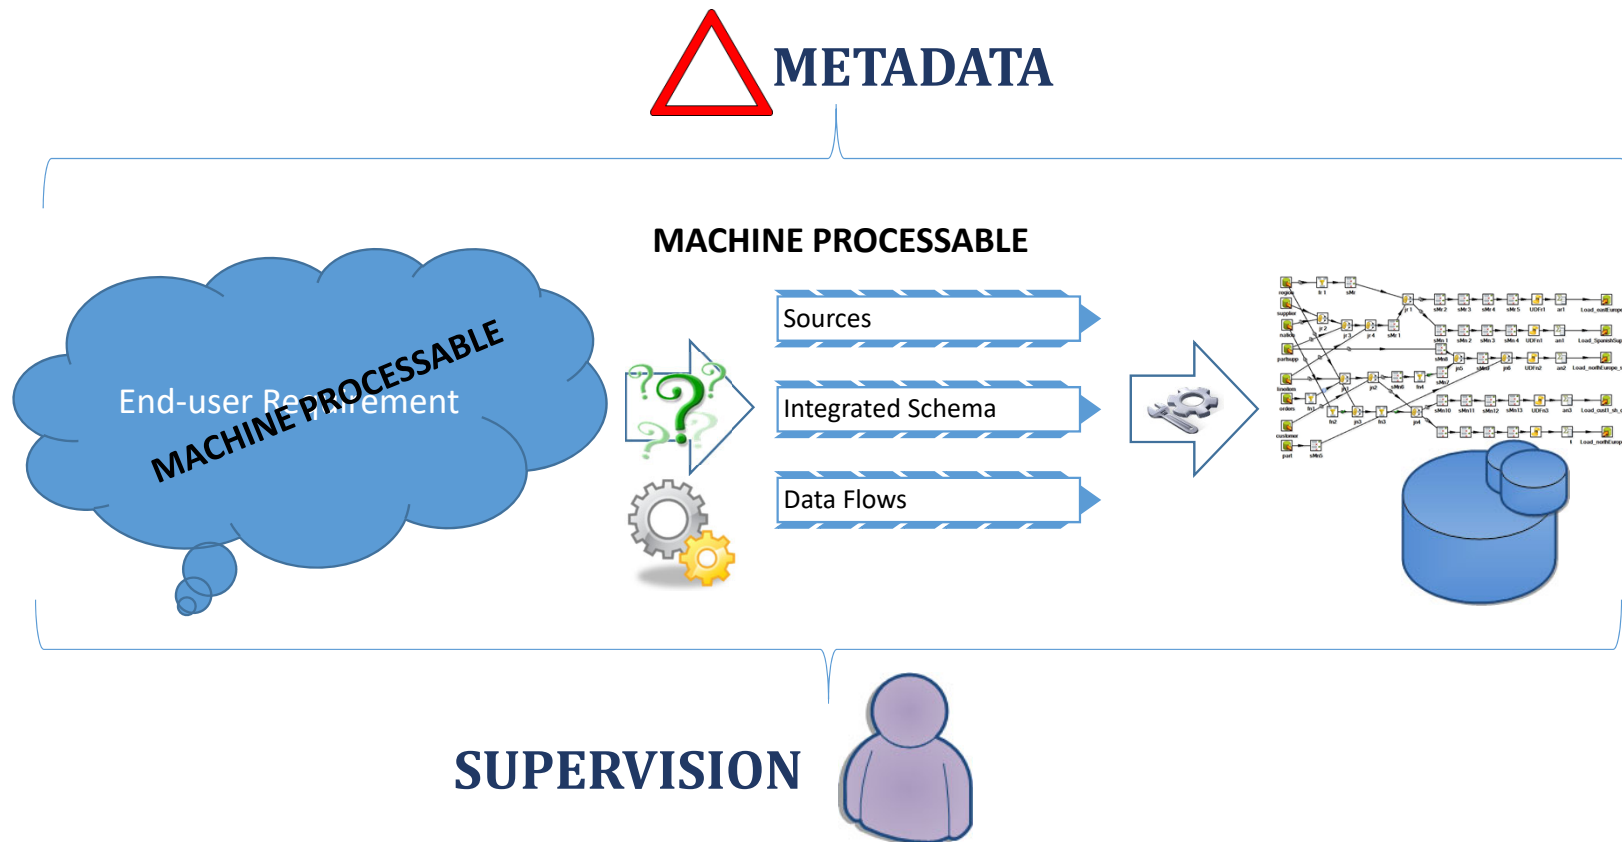
- Metadata provides semantics
  - Source schemata, mappings to views, parsing information, ...
- Automation of the integration processes

# From IT-Centered to User-Centered

Semantic-aware data repository

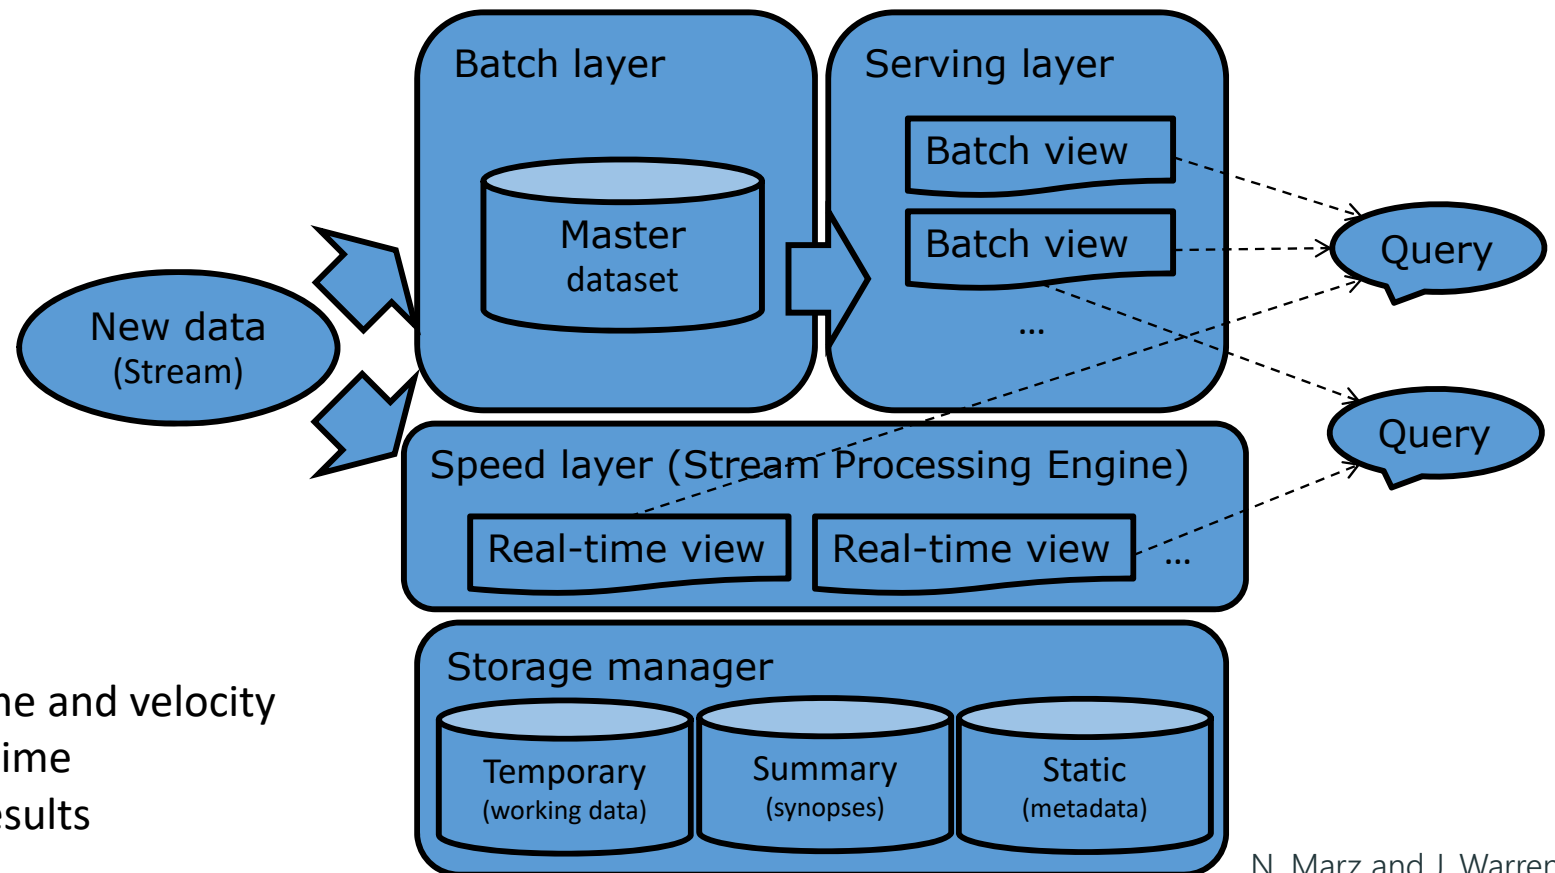


# The Missing Link: Metadata



# New Processing Architectural Patterns

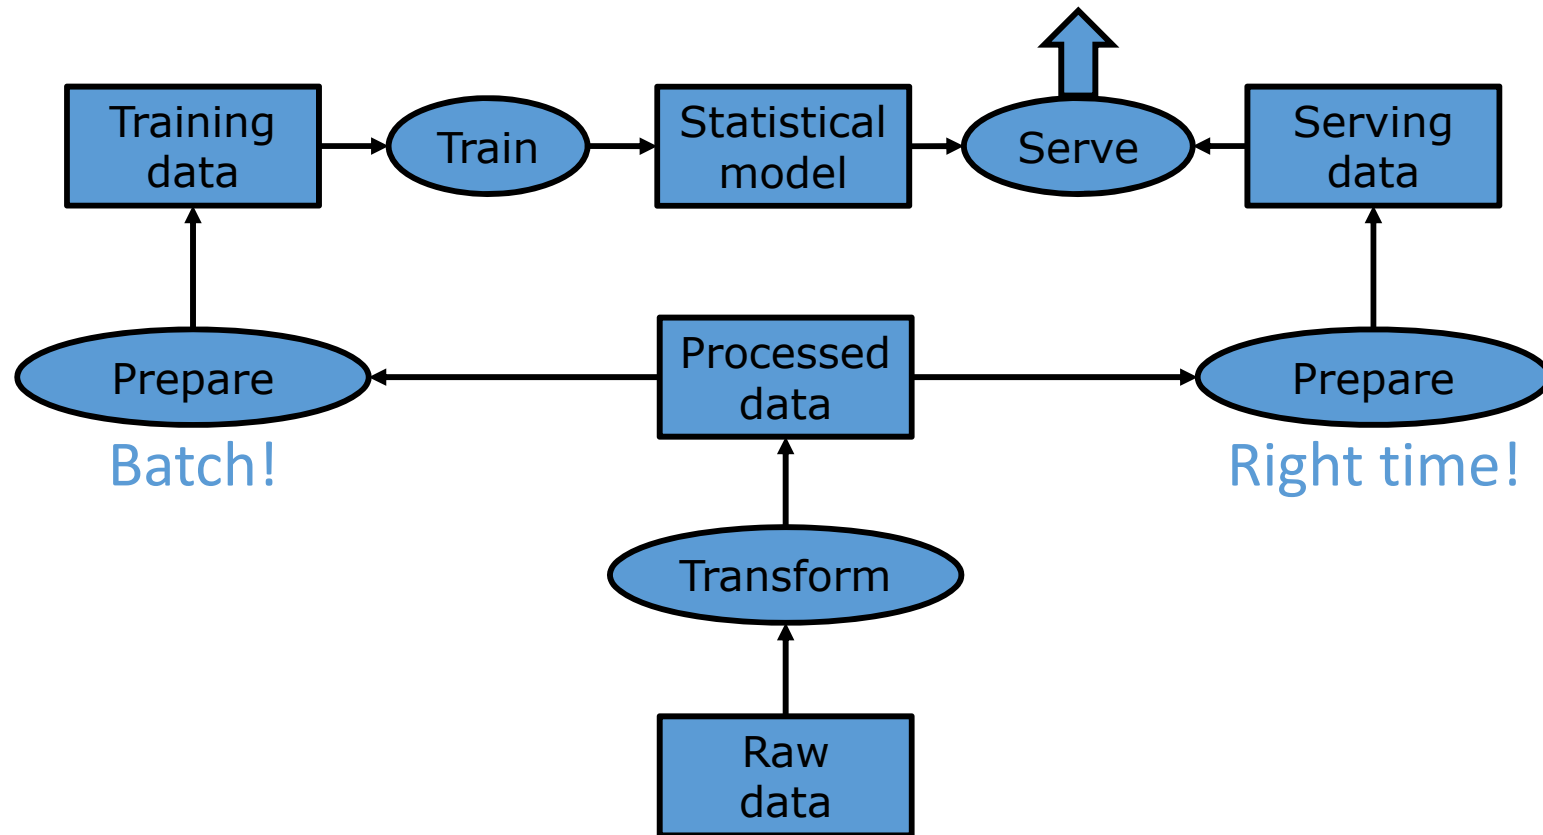
# λ-Architecture




**Idea:** Accommodate volume and velocity  
Batch processing vs. Real time  
Precise vs. Approximate results

N. Marz and J. Warren

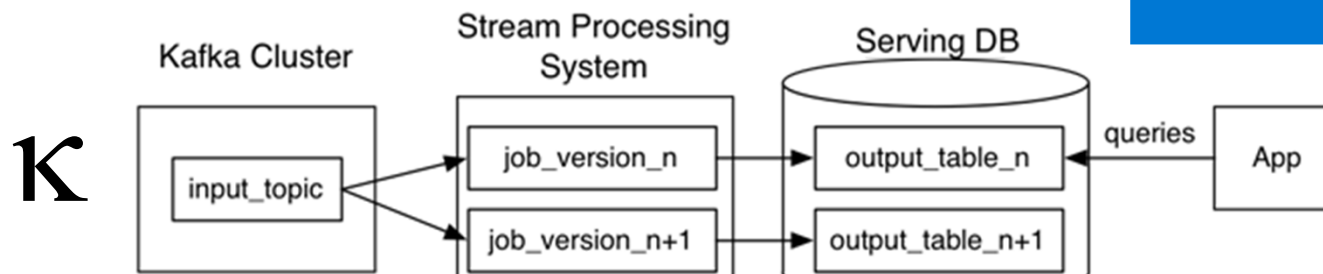
# Data-centered architecture



# $\kappa$ -architecture



- Data is considered to be a never-ending stream

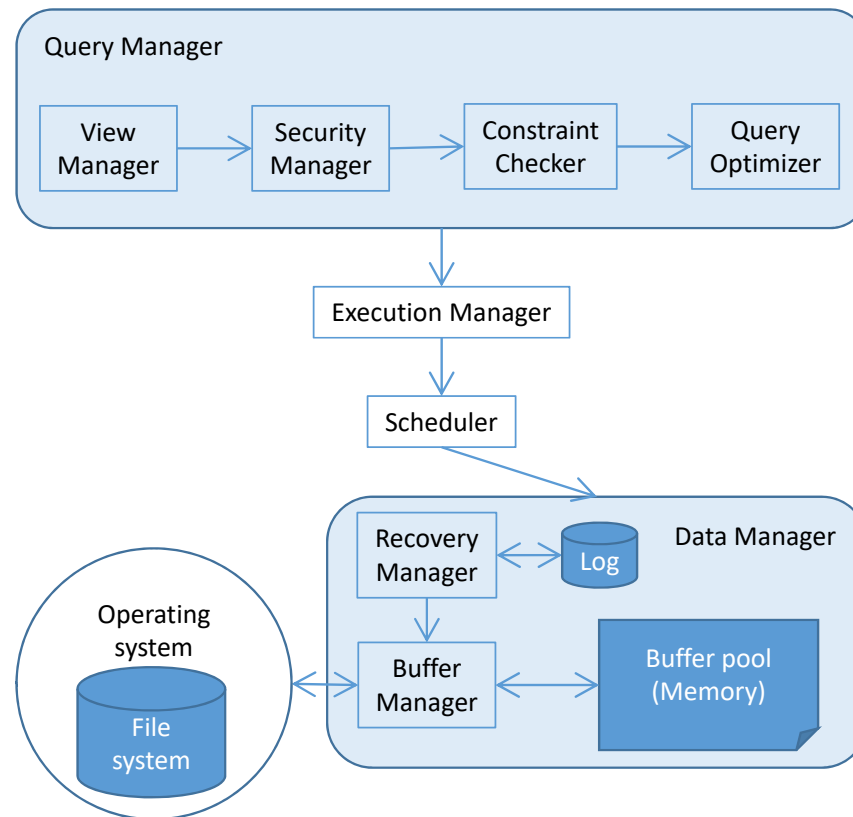


In a  $\kappa$ -architecture, data is stored as an append

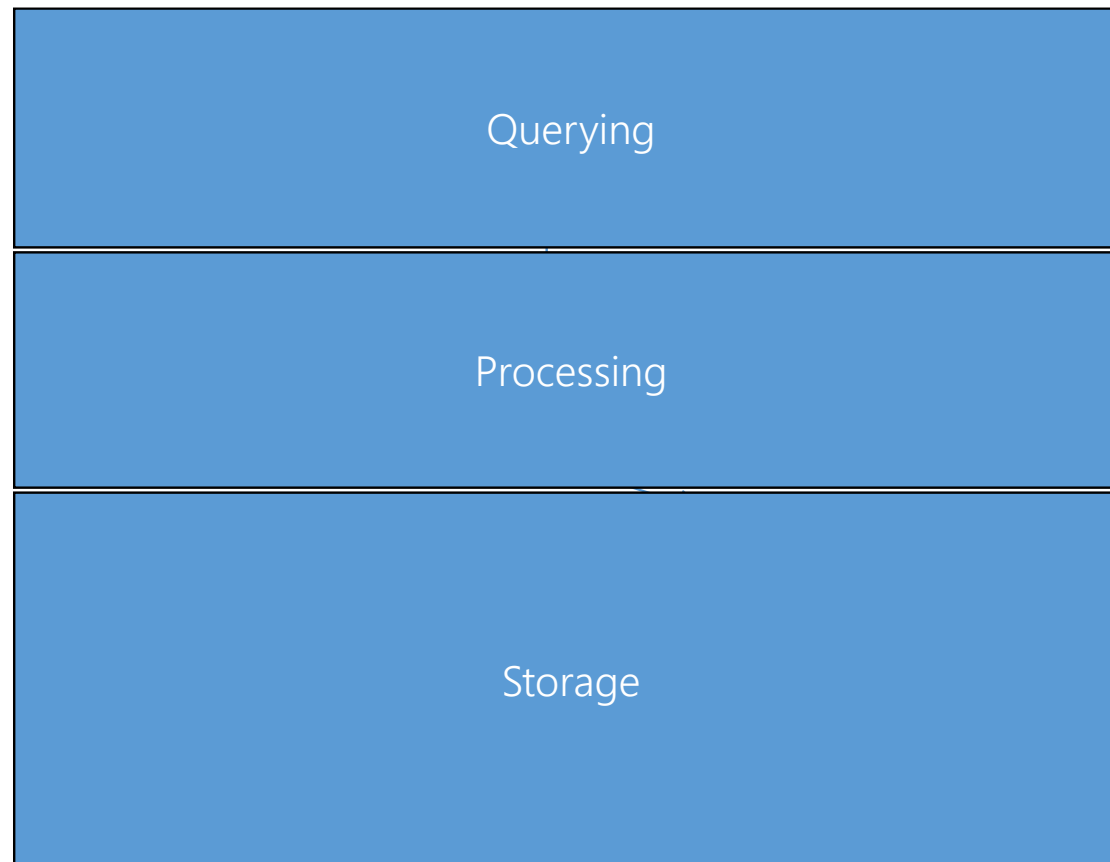


# Database Management System view

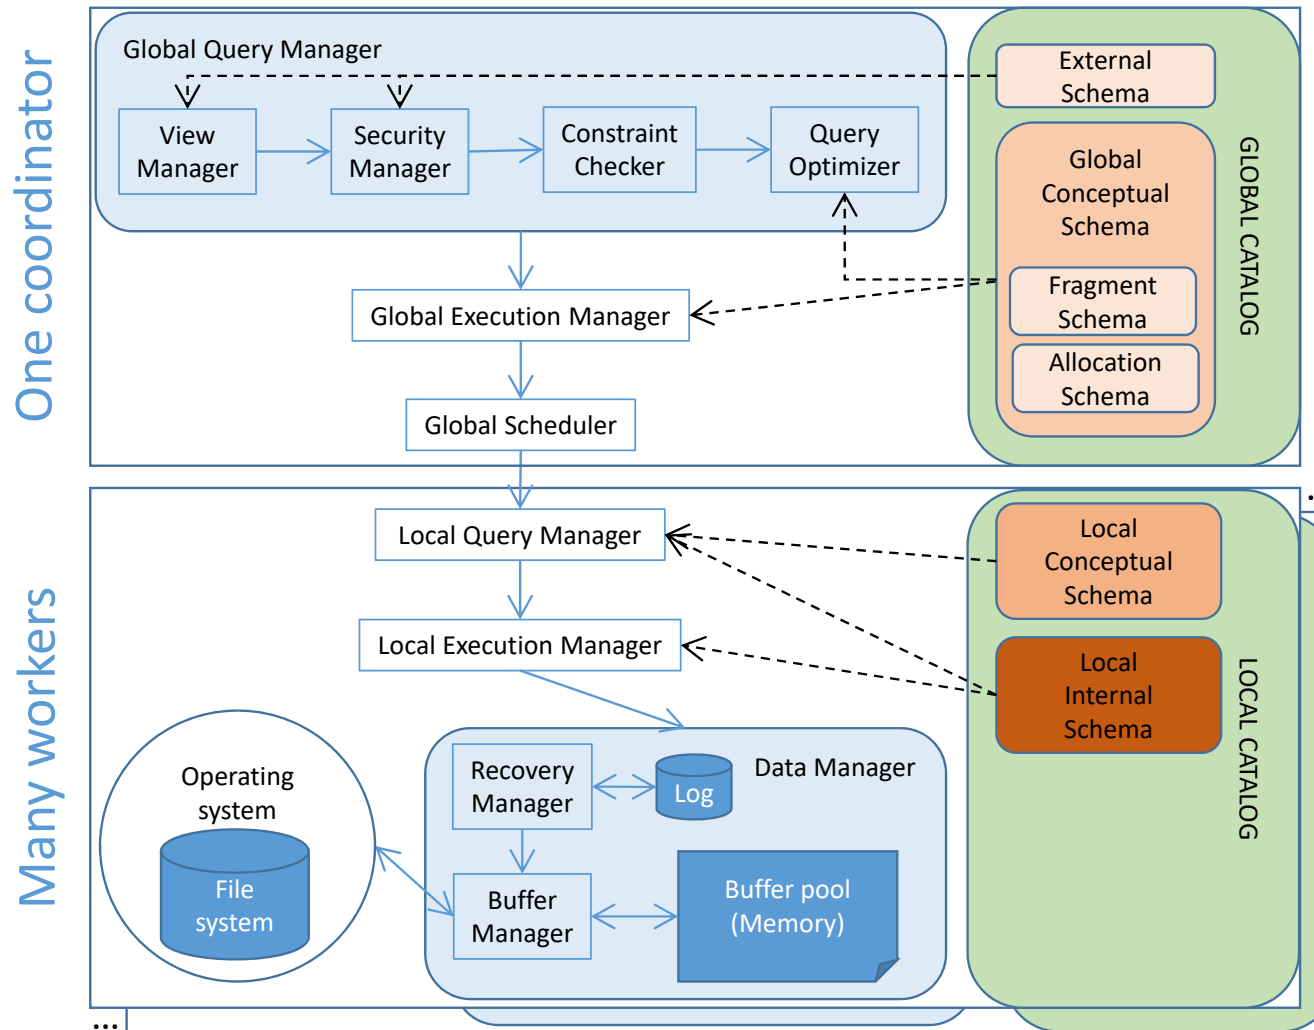
# Centralized DBMS Architecture



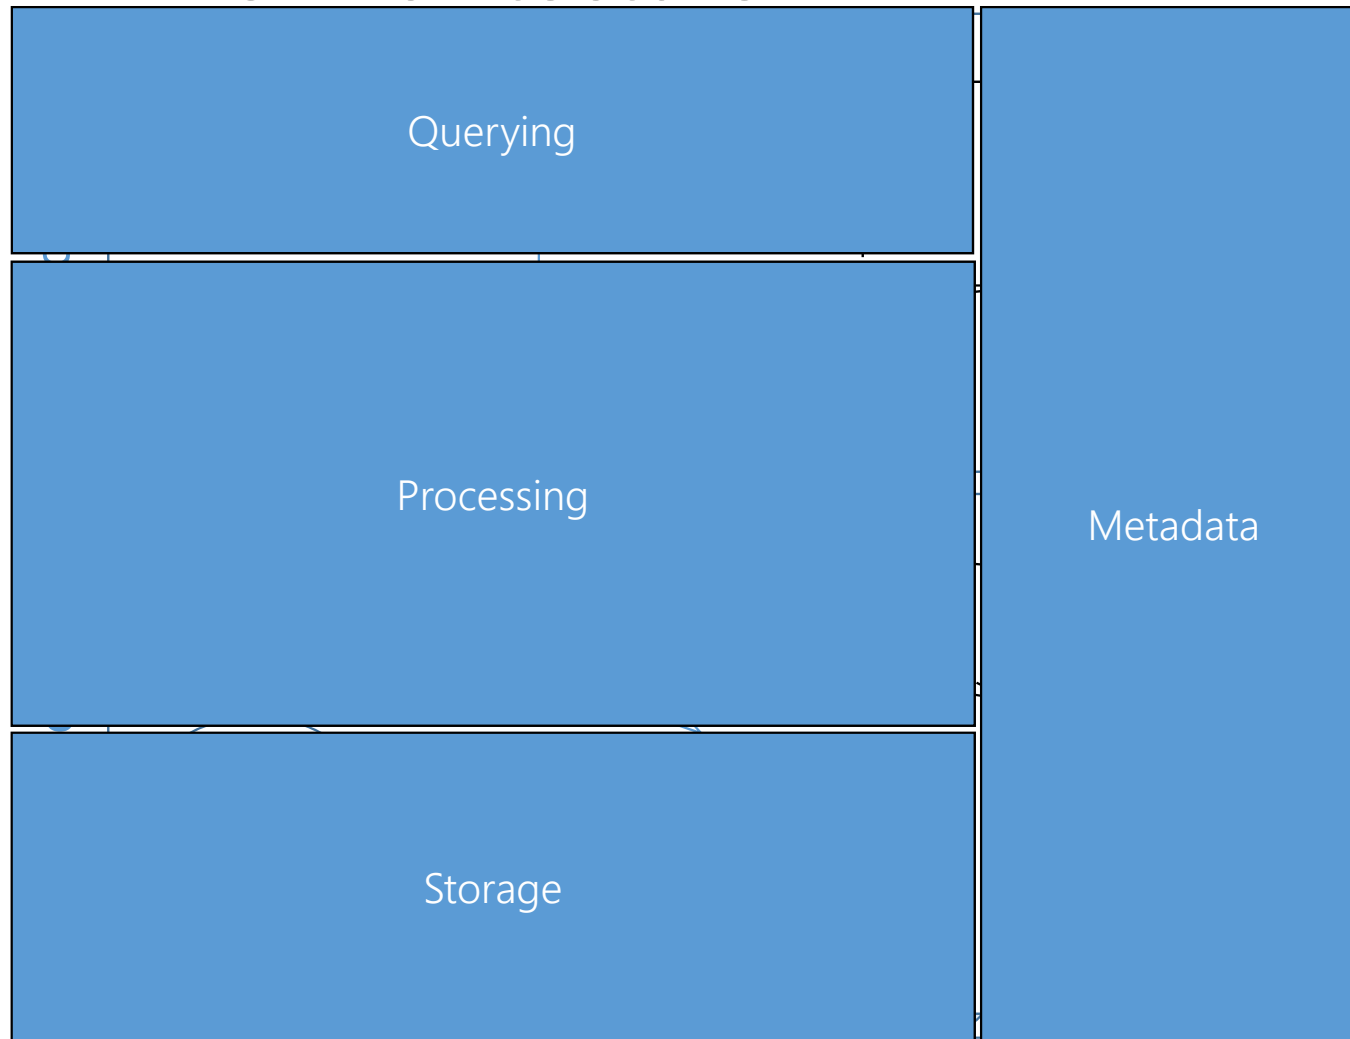
# Centralized DBMS Architecture



# Distributed DBMS Architecture




# Distributed DBMS Architecture

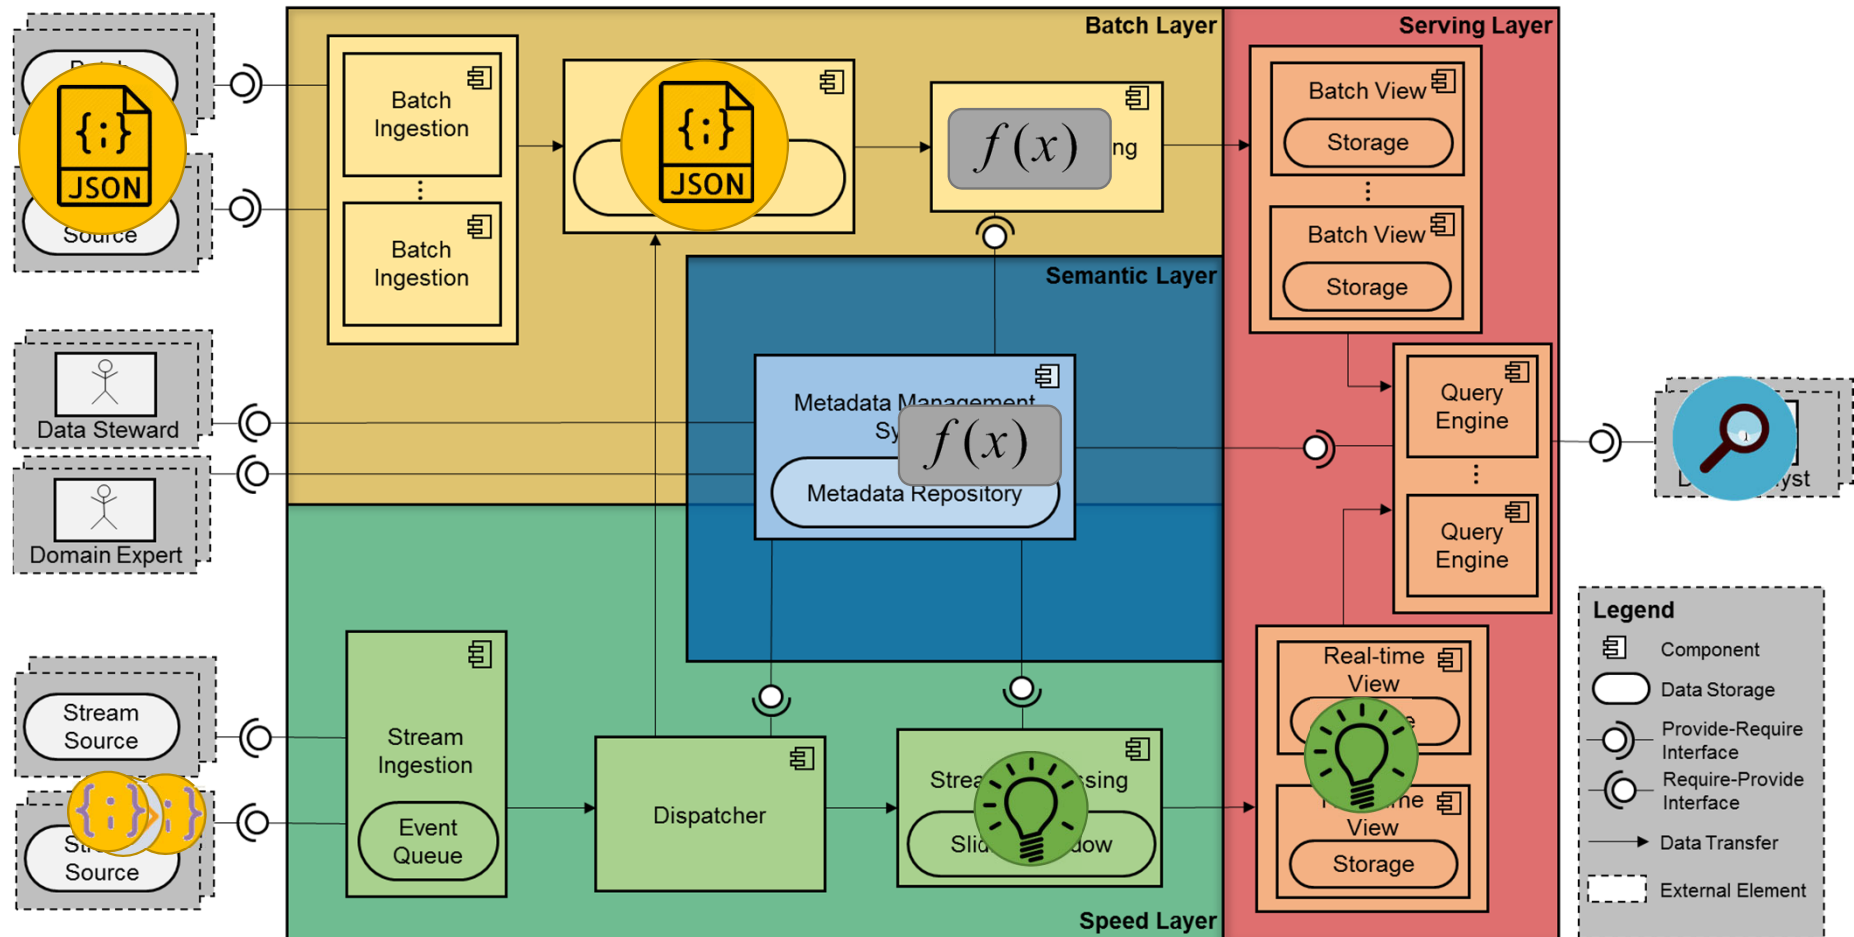


# Bolster


In batch processing, Data Lake is for storage



# Bolster



# Bolster Instantiation






# Workflow Orchestrators

Oozie is simple  
Musketeer is more complex

- Current workflow orchestrators are rather poor: Oozie
- But there are attempts for smarter approaches: the ideas behind **Musketeer** deserve special attention

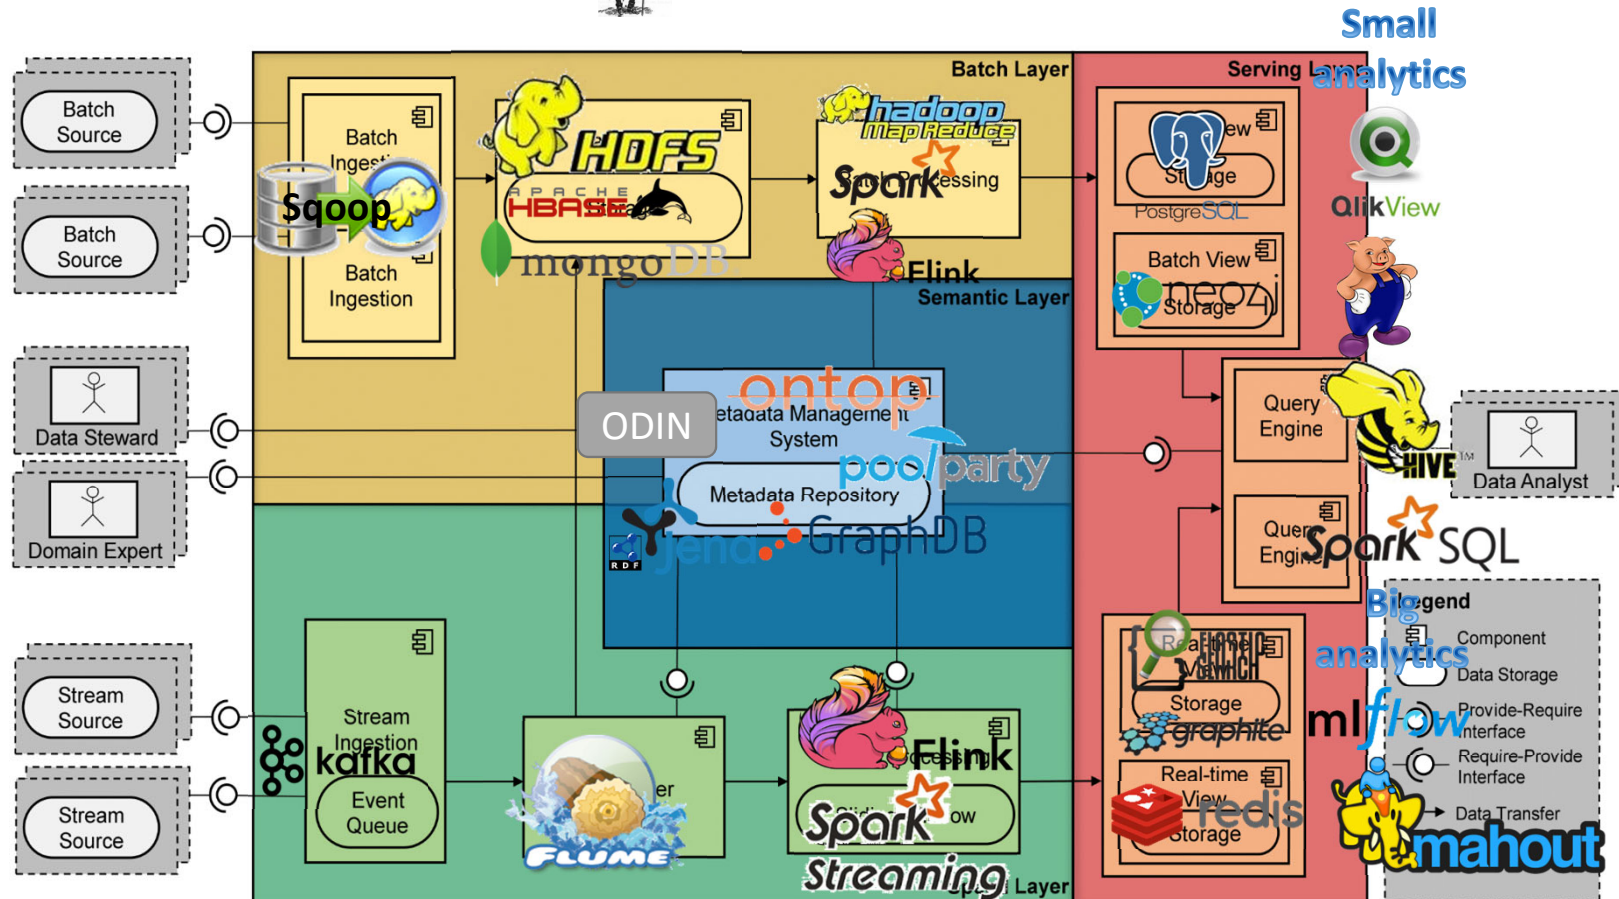


Does a similar job to  
global execution manager  
of traditional distributed  
RDBMS

# Bolster Instantiation

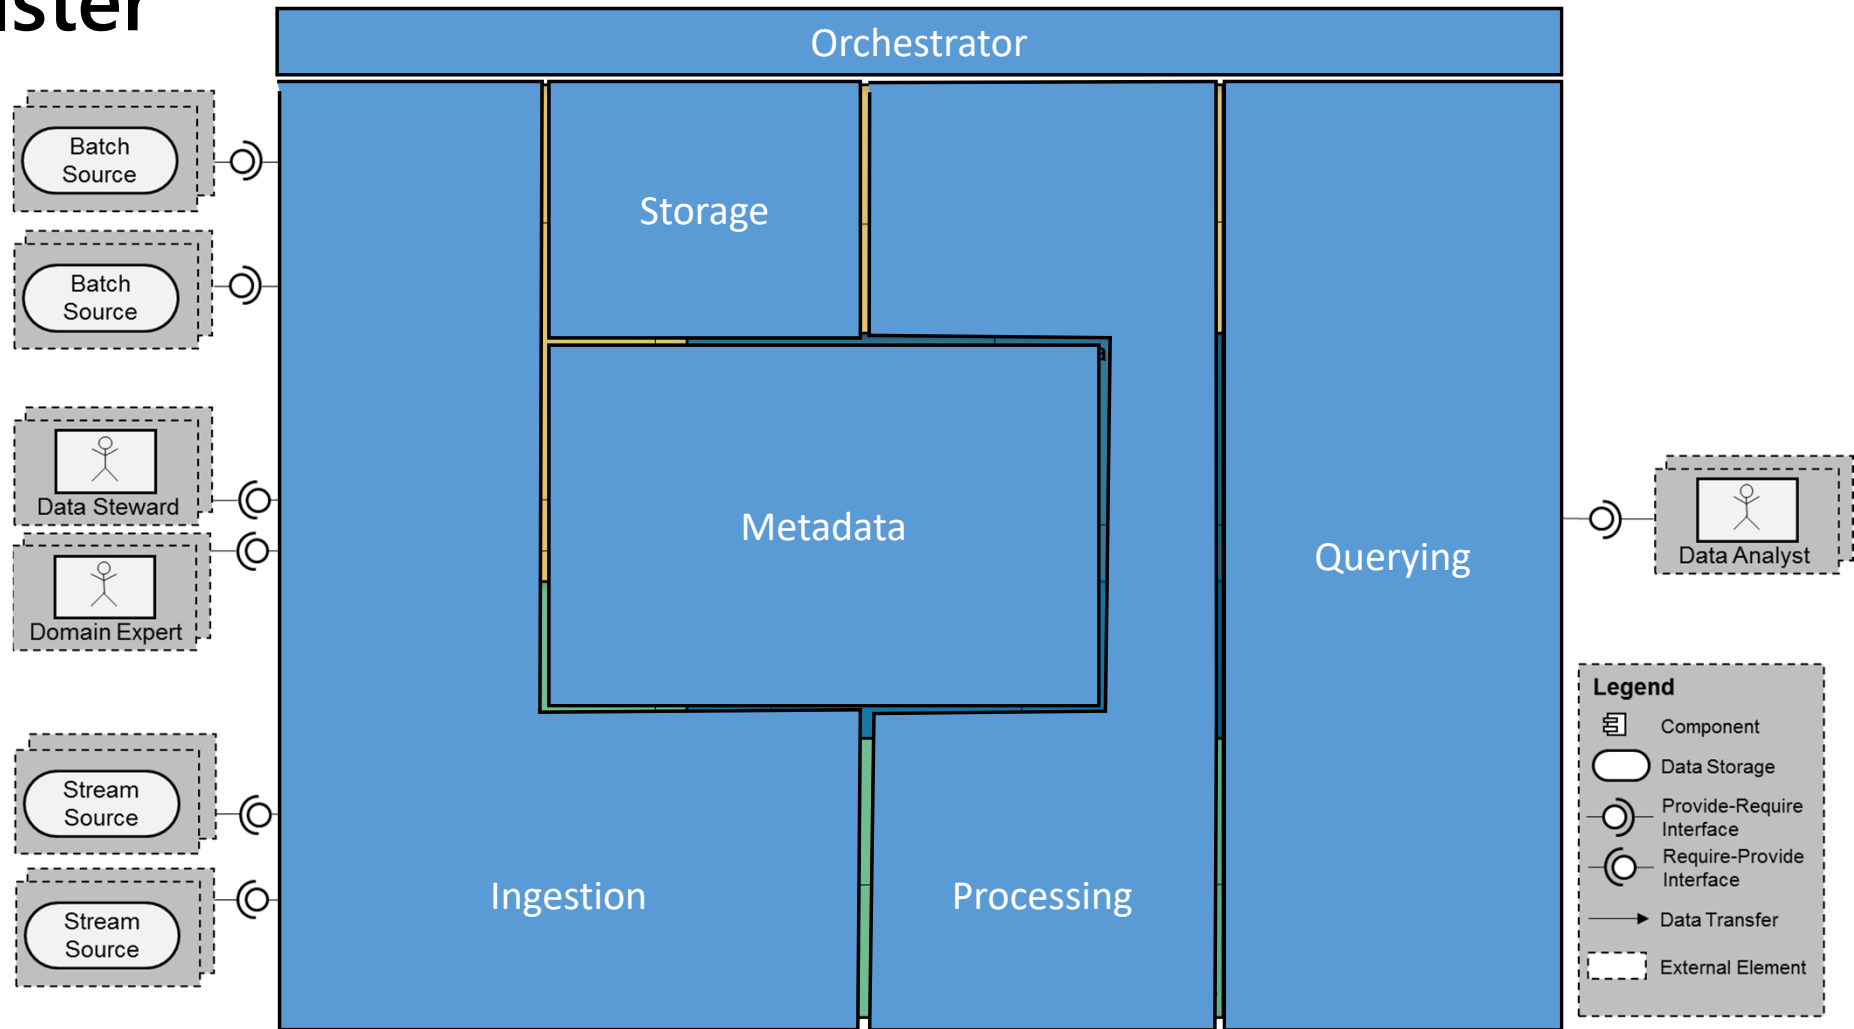


Ingestion - Sqoop, Kafka, Flume  
 Storage - HDFS+HBASE, mongodb  
 Processing - Spark, Flink, Spark Streaming, hadoop  
 metadata - hcatalog, ontop, graphdb, poolparty, jeena  
 Querying - Postgres, SparkSQL, Neo4j, redis, graphite  
 Analytics - QlikView, HIVE, SparkSQL, Big analytics, ml flow, mahout



S. Nadal et al.

# Bolster



# Closing

# Summary

- New architectural solutions
  - Lambda
  - Kappa
  - Polyglot systems
- Data Lakes
  - The need of metadata
- Reference architectures
  - Bolster
  - Quarry

# References

- D. McCreary and A. Kelly. *Making Sense of NoSQL*. Manning, 2014
- M. Grover et al. *Hadoop Application Architectures*. O'Reilly, 2015
- N. Marz and J. Warren. *Big Data: Principles and best practices of scalable realtime data systems*. Manning Publications Co., 2015
- D. Sculley et al. Hidden technical debt in machine learning systems. In *Advances in Neural Information Processing Systems*, 28. Curran Associates, Inc., 2015.
- S. Nadal et al. *A Software Reference Architecture for Semantic-Aware Big Data Systems*. Information and Software Technology 90. Elsevier, 2017
- S. Nadal et al. *ODIN: A Dataspace Management System*. International Semantic Web Conference 2019
- S. Nadal. *Metadata-Driven Data Integration* (PhD Thesis). 2019
- R. Hai, et al. *Data lake concept and systems: a survey*. CoRR abs/2106.09592. 2021
- P. Jovanovic et al. *Quarry: A User-centered Big Data Integration Platform*. Information Systems Frontier, 2021

# Resources

- <http://hadoop.apache.org>
- <http://www.cloudera.com>
- <http://hortonworks.com>